

NAME

vsearch — a versatile open-source tool for microbiome analysis, including chimera detection, clustering, dereplication and rereplication, extraction, FASTA/FASTQ/SFF file processing, masking, orienting, pairwise alignment, restriction site cutting, searching, shuffling, sorting, subsampling, and taxonomic classification of amplicon sequences for metagenomics, genomics, and population genetics.

SYNOPSIS

Chimera detection:

```
vsearch (--uchime_denovo | --uchime2_denovo | --uchime3_denovo) fastafile (--chimeras | --nonchimeras | --uchimealns | --uchimeout) outputfile [options]
vsearch --uchime_ref fastafile (--chimeras | --nonchimeras | --uchimealns | --uchimeout) outputfile --db fastafile [options]
```

Clustering:

```
vsearch (--cluster_fast | --cluster_size | --cluster_smallmem | --cluster_unequal) fastafile (--alnout | --biomout | --blast6out | --centroids | --clusters | --mothur_shared_out | --msaout | --otutabout | --profile | --samout | --uc | --userout) outputfile --id real [options]
```

Dereplication and rereplication:

```
vsearch --fastx_uniques (fastafile | fastqfile) (--fastaout | --fastqout | --tabbedout | --uc) outputfile [options]
vsearch (--derep_fulllength | --derep_id | --derep_prefix) fastafile (--output | --uc) outputfile [options]
vsearch --derep_smallmem (fastafile | fastqfile) --fastaout outputfile [options]
vsearch --rereplicate fastafile --output outputfile [options]
```

Extraction of sequences:

```
vsearch --fastx_getseq fastafile (--fastaout | --fastqout | --notmatched | --notmatchedfq) outputfile --label label [options]
vsearch --fastx_getsegs fastafile (--fastaout | --fastqout | --notmatched | --notmatchedfq) outputfile (--label label --labels labelfile | --label_word label | --label_words labelfile) [options]
vsearch --fastx_getsubseq fastafile (--fastaout | --fastqout | --notmatched | --notmatchedfq) outputfile --label label [--subseq_start position] [--subseq_end position] [options]
```

FASTA/FASTQ/SFF file processing:

```
vsearch --fasta2fastq fastafile --fastqout outputfile [options]
vsearch --fastq_chars fastqfile [options]
vsearch --fastq_convert fastqfile --fastqout outputfile [options]
vsearch (--fastq_eestats | --fastq_eestats2) fastqfile --output outputfile [options]
vsearch --fastq_filter fastqfile [--reverse fastqfile] (--fastaout | --fastaout_discarded | --fastqout | --fastqout_discarded --fastaout_rev | --fastaout_discarded_rev | --fastqout_rev | --fastqout_discarded_rev) outputfile [options]
vsearch --fastq_join fastqfile --reverse fastqfile (--fastaout | --fastqout) outputfile [options]
vsearch --fastq_mergepairs fastqfile --reverse fastqfile (--fastaout | --fastqout | --fastaout_notmerged_fwd | --fastaout_notmerged_rev | --fastqout_notmerged_fwd | --fastqout_notmerged_rev | --eetabbedout) outputfile [options]
vsearch --fastq_stats fastqfile [-log logfile] [options]
vsearch --fastx_filter inputfile [--reverse inputfile] (--fastaout | --fastaout_discarded | --fastqout | --fastqout_discarded --fastaout_rev | --fastaout_discarded_rev | --fastqout_rev | --fastqout_discarded_rev) outputfile [options]
vsearch --fastx_revcomp inputfile (--fastaout | --fastqout) outputfile [options]
```

vsearch --sff_convert *sff-file* --fastqout *outputfile* [*options*]

Masking:

vsearch --fastx_mask *fastxfile* (--fastaout | --fastqout) *outputfile* [options]

vsearch --maskfasta *fastafilename* --output *outputfile* [*options*]

Orienting:

vsearch --orient *fastxfile* --db *fastxfile* (--fastaout | --fastqout | --notmatched | --tabbedout) *output-file [options]*

Pairwise alignment:

```
vsearch --allpairs_global fastafile (--alnout | --blast6out | --matched | --notmatched | --samout |  
--uc | --userout) outputfile (--acceptall | --id real) [options]
```

Restriction site cutting:

vsearch --cut *fastafilename* --cut_pattern *pattern* (--fastaout | --fastaout_rev | --fastaout_discarded | --fastaout_discarded_rev) *outputfile* [options]

Searching:

vsearch --search_exact *fastafilename* --db *fastafilename* (--alnout | --biomout | --blast6out | --mothur shared out | --otutabout | --samout | --uc | --userout | --lcaout) *outputfile* [options]

```
vsearch --usearch_global fastafile --db (fastafile | udbfile) (--alnout | --biomout | --blast6out | --mothur_shared_out | --otutabout | --samout | --uc | --userout | --lcaout) outputfile --id real [options]
```

Shuffling and sorting:

vsearch (**--shuffle** | **--sortbylength** | **--sortbysize**) *fastafile* --output *outputfile* [*options*]

Subsampling:

vsearch --fastx_subsample *fastafilename* (--fastaout | --fastqout) *outputfile* (--sample_pct *real* | --sample_size *positive integer*) [*options*]

Taxonomic classification:

vsearch --syntax *fastafile* --db (*fastafile* | *udbfile*) --tabbedout *outputfile* [--syntax_cutoff *real*] [*options*]

UDB database handling:

```
vsearch --makeudb_usearch fastafilename --output outputfile [options]
```

vsearch --fdb2fasta *fdbfile* --output *outputfile* [*options*]

vsearch (**--fdbinfo** | **--fdbstats**) *fdbfile* [*options*]

DESCRIPTION

Environmental or clinical molecular diversity studies generate large volumes of amplicons (e.g.; SSU-rRNA sequences) that need to be checked for chimeras, dereplicated, masked, sorted, searched, clustered or compared to reference sequences. The aim of **vsearch** is to offer a all-in-one open source tool to perform these tasks, using optimized algorithm implementations and harvesting the full potential of modern computers, thus providing fast and accurate data processing.

Comparing nucleotide sequences is at the core of **vsearch**. To speed up comparisons, **vsearch** implements an extremely fast Needleman-Wunsch algorithm, making use of the Streaming SIMD Extensions (SSE2) of post-2003 x86-64 CPUs. If SSE2 instructions are not available, **vsearch** exits with an error message. On Power8 CPUs it will use AltiVec/VSX/VMX instructions, and on ARMv8 CPUs it will use Neon instructions. On other systems it can use the SIMD Everywhere (simde) library, if available. Memory usage increases rapidly with sequence length: for example comparing two sequences of length 1 kb requires 8 MB of memory per thread, and comparing two 10 kb sequences requires 800 MB of memory per thread. For comparisons involving sequences with a length product greater than 25 million (for example two sequences of length 5 kb), **vsearch** uses a slower alignment method described by Hirschberg (1975) and Myers and Miller (1988), with much smaller memory requirements.

Input

vsearch accept as input fasta or fastq files containing one or several nucleotidic entries. In fasta files, each entry is made of a header and a sequence. The header is defined as the string comprised between the initial '`>`' symbol and the first space, tab or the end of the line, unless the `--notrunclabels` option is in effect, in which case the entire line is included. The header should contain printable ascii characters (33-126). The program will terminate with a fatal error if there are unprintable ascii characters. A warning will be issued if non-ascii characters (128-255) are encountered.

If the header matches the pattern '`>[;]size=integer;label`', the pattern '`>label;size=integer;label`', or the pattern '`>label;size=integer[;]`', **vsearch** will interpret `integer` as the number of occurrences (or abundance) of the sequence in the study. That abundance information is used or created during chimera detection, clustering, dereplication, sorting and searching.

The sequence is defined as a string of IUPAC symbols (ACGTURYSWKMDBHVN), starting after the end of the identifier line and ending before the next identifier line, or the file end. **vsearch** silently ignores ascii characters 9 to 13, and exits with an error message if ascii characters 0 to 8, 14 to 31, '.' or '-' are present. All other ascii or non-ascii characters are stripped and complained about in a warning message.

In fastq files, each entry is made of sequence header starting with a symbol '@', a nucleotidic sequence (same rules as for fasta sequences), a quality header starting with a symbol '+' and a string of ASCII characters (offset 33 or 64), each one encoding the quality value of the corresponding position in the nucleotidic sequence.

vsearch operations are case insensitive, except when soft masking is activated. DUST masking is automatically applied during chimera detection, clustering, masking, pairwise alignment and searching. Soft masking is specified with the options '`--dbmask soft`' (for searching and chimera detection with a reference) or '`--qmask soft`' (for searching, *de novo* chimera detection, clustering and masking). When using soft masking, lower case letters indicate masked symbols, while upper case letters indicate regular symbols. Masked symbols are never included in the unique index words used for sequence comparisons, otherwise they are treated as normal symbols.

When comparing sequences during chimera detection, dereplication, searching and clustering, T and U are considered identical, regardless of their case. When aligning sequences, identical symbols will receive a positive match score (default +2). If two symbols are not identical, their alignment result in a negative mismatch score (default -4). Aligning a pair of symbols where at least one of them is an ambiguous symbol (BDHKMNRSVWY) will always result in a score of zero. Alignment of two identical ambiguous symbols (for example, R vs R) also receives a score of zero. When computing the amount of similarity by counting matches and mismatches after alignment, ambiguous nucleotide symbols will count as matching to other symbols if they have at least one of the nucleotides (ACGTU) they may represent in common. For example: W will match A and T, but also any of MRVHDN. When showing alignments (for example with the `--alnout` option) matches involving ambiguous symbols will be shown with a plus character (+) between them while exact matches between non-ambiguous symbols will be shown with a vertical bar character (|).

vsearch can read data from standard files and write to standard files, but it can also read from pipes and write to pipes! For example, multiple fasta files can be piped into **vsearch** for dereplication. To do so, file names can be replaced with:

- the symbol '`-`', representing '`/dev/stdin`' for input files or '`/dev/stdout`' for output files (with an exception for '`--db -`', see * below),
- a named pipe created with the command `mkfifo`,
- a process substitution '`<(command)`' as input or '`>(command)`' as output.

* `--db -` is not accepted, to prevent potential concurrent reads from `stdin`. A workaround for advanced users is to call '`--db /dev/stdin`' directly.

vsearch can automatically read compressed gzip or bzip2 files if the appropriate libraries are present during the compilation. **vsearch** can also read pipes streaming compressed gzip or bzip2 data if the options `--gzip_decompress` or `--bzip2_decompress` are selected. When reading from a pipe, the progress indicator is not updated.

Options

vsearch recognizes a large number of command-line commands and options. For easier navigation, options are grouped below by theme (chimera detection, clustering, dereplication and rereplication, FASTA/FASTQ file processing, masking, pairwise alignment, searching, shuffling, sorting, and subsampling). We start with the general options that apply to all themes. Options start with a double dash (--). A single dash (-) may also be used, except on NetBSD systems. Option names may be shortened as long as they are not ambiguous (e.g. --derep_f).

Help and version commands:

--help -h Display help text with brief information about all commands and options.

--version -v

Output version information and a citation for the VSEARCH publication. Show the status of the support for gzip- and bzip2-compressed input files.

General options:

--bzip2_decompress

When reading from a pipe streaming bzip2-compressed data, decompress the data. This option is not needed when reading from a standard bzip2-compressed file.

--fasta_width positive integer

Fasta files produced by **vsearch** are wrapped (sequences are written on lines of *integer* nucleotides, 80 by default). Set the value to zero to eliminate the wrapping.

--gzip_decompress

When reading from a pipe streaming gzip-compressed data, decompress the data. This option is not needed when reading from a standard gzip-compressed file.

--label_suffix string

When writing FASTA or FASTQ files, add the suffix *string* to sequence headers.

--log filename

Write messages to the specified log file. Information written includes program version, amount of memory available, number of cores and command line options, and if need be, informational messages, warnings and fatal errors. The start and finish times are also recorded as well as the elapsed time and the maximum amount of memory consumed. The different **vsearch** commands can also write additional information to the log file.

--maxseqlength positive integer

All **vsearch** operations discard sequences longer than *integer* (50,000 nucleotides by default).

--minseqlength positive integer

All **vsearch** operations discard sequences shorter than *integer*: 1 nucleotide by default for sorting or shuffling, 32 nucleotides for clustering and dereplication as well as the commands --makeudb_usearch, --sintax, and --usearch_global.

--no_progress

Do not show the gradually increasing progress indicator.

--notrunclabels

Do not truncate sequence labels at first space or tab, but use the full header in output files. Turned off by default for all commands except the sintax command.

--quiet Suppress all messages to stdout and stderr except for warnings and fatal error messages.

--sample string

When writing FASTA or FASTQ files, add the the given sample identifier *string* to sequence headers. For instance, if the given string is ABC, the text ";sample=ABC" will

be added to the header. Note that *string* will be truncated at the first ';' or blank character. Other characters (alphabetical, numerical and punctuations) are accepted.

--threads *positive integer*

Number of computation threads to use (1 to 1024). The number of threads should be less than or equal to the number of available CPU cores. The default is to use all available resources and to launch one thread per core. The following commands are multi-threaded: allpairs_global, cluster_fast, cluster_size, cluster_smallmem, cluster_ounoise, fastq_mergepairs, fastx_mask, maskfasta, search_exact, syntax, uchime_ref, and usearch_global. Only one thread is used for the other commands.

Chimera detection options:

Chimera detection is based on a scoring function controlled by five options (--dn, --mindiffs, --mindiv, --minh, --xn). Sequences are first sorted by decreasing abundance, if available, and compared on their *plus* strand only (case insensitive).

Input sequences are masked as specified with the --qmask and --hardmask options. Masking of the database for reference based chimera detection is specified with the --dbmask option.

In *de novo* mode, input fasta file must present abundance annotations (i.e. a pattern [:]size=integer[:] in the fasta header). Input order matters for chimera detection, so we recommend to sort sequences by decreasing abundance (default of --derep_fulllength command). If your sequence set needs to be sorted, please see the --sortbysize command in the sorting section.

--abskew *real*

When using --uchime_denovo, the abundance skew is used to distinguish in a three-way alignment which sequence is the chimera and which are the parents. The assumption is that chimeras appear later in the PCR amplification process and are therefore less abundant than their parents. For --uchime3_denovo the default value is 16.0. For the other commands, the default value is 2.0, which means that the parents should be at least 2 times more abundant than their chimera. Any positive value equal or greater than 1.0 can be used.

--alignwidth *positive integer*

When using --uchimealns, set the width of the three-way alignments (80 nucleotides by default). Set to zero to eliminate wrapping.

--borderline *filename*

Output borderline chimeric sequences to *filename*, in fasta format. Borderline chimeric sequences are sequences that have a high enough score but which are not sufficiently different from their closest parent.

--chimeras *filename*

Output chimeric sequences to *filename*, in fasta format. Output order may vary when using multiple threads.

--db *filename*

When using --uchime_ref, detect chimeras using the reference sequences contained in *filename*. Reference sequences are assumed to be chimera-free. Chimeras cannot be detected if their parents, or sufficiently close relatives, are not present in the database. The file name must refer to a FASTA file or to a UDB file. If a UDB file is used, it should be created using the --makeudb_usearch command with the --dbmask dust option.

--dn *strictly positive real number*

pseudo-count prior on the number of no votes, corresponding to the parameter *n* in the chimera scoring function (default value is 1.4). Increasing --dn reduces the likelihood of tagging a sequence as a chimera (less false positives, but also more false negatives).

--fasta_score

Add the chimera score to the headers in the fasta output files for chimeras, non-chimeras and borderline sequences, using the format ';uchime_denovo=*float*'.

--lengthout

Write sequence length information to the output files in FASTA format by adding a ";length=*integer*" attribute in the header.

--mindiffs *positive integer*

Minimum number of differences per segment (default value is 3). The parameter is ignored with --uchime2_denovo and --uchime3_denovo.

--mindiv *real*

Minimum divergence from closest parent (default value is 0.8). The parameter is ignored with --uchime2_denovo and --uchime3_denovo.

--minh *real*

Minimum score (*h*). Increasing this value tends to reduce the number of false positives and to decrease sensitivity. Default value is 0.28, and values ranging from 0.0 to 1.0 included are accepted. The parameter is ignored with --uchime2_denovo and --uchime3_denovo.

--nonchimeras *filename*

Output non-chimeric sequences to *filename*, in fasta format. Output order may vary when using multiple threads.

--relabel *string*

Relabel sequences using the prefix *string* and a ticker (1, 2, 3, etc.) to construct the new headers. Use --sizeout to conserve the abundance annotations.

--relabel_keep

When relabelling, keep the old identifier in the header after a space.

--relabel_md5

Relabel sequences using the MD5 message digest algorithm applied to each sequence. Former sequence headers are discarded. The sequence is converted to upper case and each 'U' is replaced by a 'T' before computation of the digest. The MD5 digest is a cryptographic hash function designed to minimize the probability that two different inputs give the same output, even for very similar, but non-identical inputs. Still, there is a very small, but non-zero, probability that two different inputs give the same digest (i.e. a collision). MD5 generates a 128-bit (16-byte) digest that is represented by 16 hexadecimal numbers (using 32 symbols among 0123456789abcdef). Use --sizeout to conserve the abundance annotations.

--relabel_self

Relabel sequences using each sequence itself as a label.

--relabel_sha1

Relabel sequences using the SHA1 message digest algorithm applied to each sequence. It is similar to the --relabel_md5 option but uses the SHA1 algorithm instead of the MD5 algorithm. SHA1 generates a 160-bit (20-byte) digest that is represented by 20 hexadecimal numbers (40 symbols). The probability of a collision (two non-identical sequences resulting in the same digest) is smaller for the SHA1 algorithm than it is for the MD5 algorithm.

--self

When using --uchime_ref, ignore a reference sequence when its label matches the label of the query sequence (useful to estimate false-positive rate in reference sequences).

--selfid

When using --uchime_ref, ignore a reference sequence when its nucleotide sequence is strictly identical to the nucleotidic sequence of the query.

--sizein

In *de novo* mode, abundance annotations (pattern '[>:]size=*integer*[;]') present in sequence headers are taken into account by default (--sizein is always implied). This option is ignored by --uchime_ref.

--sizeout When relabelling, add abundance annotations to fasta headers (using the format '`;size=integer;`').

--uchime_denovo *filename*

Detect chimeras present in the fasta-formatted *filename*, without external references (i.e. *de novo*). Automatically sort the sequences in *filename* by decreasing abundance beforehand (see the sorting section for details). Multithreading is not supported.

--uchime2_denovo *filename*

Detect chimeras present in the fasta-formatted *filename*, using the UCHIME2 algorithm. This algorithm is designed for denoised amplicons (see --cluster_unequal). Automatically sort the sequences in *filename* by decreasing abundance beforehand (see the sorting section for details). Multithreading is not supported.

--uchime3_denovo *filename*

Detect chimeras present in the fasta-formatted *filename*, using the UCHIME2 algorithm. The only difference from --uchime2_denovo is that the default minimum abundance skew (--abskew) is set to 16.0 rather than 2.0.

--uchime_ref *filename*

Detect chimeras present in the fasta-formatted *filename* by comparing them with reference sequences (option --db). Multithreading is supported.

--uchimealns *filename*

Write the three-way global alignments (parentA, parentB, chimera) to *filename* using a human-readable format. Use --alignwidth to modify alignment length. Output order may vary when using multiple threads. All sequences are converted to upper case before alignment. Lower case letters indicate disagreement in the alignment.

--uchimeout *filename*

Write chimera detection results to *filename* using a 18-field, tab-separated uchime-like format. Use --uchimeout5 to use a format compatible with usearch v5 and earlier versions. Rows output order may vary when using multiple threads.

1. score: higher score means a more likely chimeric alignment.
2. Q: query sequence label.
3. A: parent A sequence label.
4. B: parent B sequence label.
5. T: top parent sequence label (i.e. parent most similar to the query). That field is removed when using --uchimeout5.
6. idQM: percentage of similarity of query (Q) and model (M) constructed as a part of parent A and a part of parent B.
7. idQA: percentage of similarity of query (Q) and parent A.
8. idQB: percentage of similarity of query (Q) and parent B.
9. idAB: percentage of similarity of parent A and parent B.
10. idQT: percentage of similarity of query (Q) and top parent (T).
11. LY: yes votes in the left part of the model.
12. LN: no votes in the left part of the model.
13. LA: abstain votes in the left part of the model.
14. RY: yes votes in the right part of the model.
15. RN: no votes in the right part of the model.

16. RA: abstain votes in the right part of the model.
17. div: divergence, defined as (idQM - idQT).
18. YN: query is chimeric (Y), or not (N), or is a borderline case (?).

--uchimeout5

When using --uchimeout, write chimera detection results using a 17-field, tab-separated uchime-like format (drop the 5th field of --uchimeout), compatible with usearch version 5 and earlier versions.

--xlength Strip header attribute ";length=*integer*" from input sequences. This attribute is added to output sequences by the --lengthout option.

--xn *strictly positive real number*

weight of no votes, corresponding to the parameter *beta* in the scoring function (default value is 8.0). Increasing --xn reduces the likelihood of tagging a sequence as a chimera (less false positives, but also more false negatives).

--xsize Strip abundance information from the headers when writing the output file.

Clustering options:

vsearch implements a single-pass, greedy centroid-based clustering algorithm, similar to the algorithms implemented in usearch, DNAclust and sumaclust for example. Important parameters are the global clustering threshold (--id) and the pairwise identity definition (--iddef).

Input sequences are masked as specified with the --qmask and --hardmask options.

--biomout *filename*

Generate an OTU table in the biom version 1.0 JSON file format as specified at ([link](https://biom-format.org/documentation/format_versions/biom-1.0.html)) <https://biom-format.org/documentation/format_versions/biom-1.0.html>. The format describes how to store a sparse matrix containing the abundances of the OTUs in the different samples. This format is much more efficient than the classic and mothur OTU table formats available with the --otutabout and --mothur_shared_out options, respectively, and is recommended at least for large tables. The OTUs are represented by the cluster centroids. Taxonomy information will be included for the OTUs if available. Sample identifiers will be extracted from the headers of all sequences in the input file. If the header contains ';sample=abc123;' or ';barcodeLabel=abc123;' or a similar string somewhere, then the given sample identifier (here 'abc123') will be used. The semicolon is not mandatory at the beginning or end of the header. The sample identifier may contain any printable character except semicolons. If no such sample label is found, the identifier in the initial part of the header will be used, but only letters, digits and underscores are allowed. OTU identifiers will be extracted from the headers of the cluster centroid sequences. If the header contains ';otu=def789;' or a similar string somewhere, then the given OTU identifier (here 'def789') will be used. The semicolon is not mandatory at the beginning or end of the header. The OTU identifier may contain any printable character except semicolons. If no such OTU label is found, the identifier in the initial part of the header will be used, and all characters except semicolons are allowed. Alternatively, OTU identifiers can be generated using the relabelling options (--relabel, --relabel_self, --relabel_sha1, or --relabel_md5). Taxonomy information, if present, will also be extracted from the headers of the centroid sequences. If the header contains ';tax=Homo_sapiens;' or a similar string somewhere, then the given taxonomy information (here 'Homo_sapiens') will be used. The semicolon is not mandatory at the beginning or end of the header. The taxonomy information may contain any printable character except semicolons. If an OTU table in the biom version 2.1 HDF5 file format is required, the biom utility may be used as described at ([link](https://biom-format.org/documentation/biom_conversion.html)) <https://biom-format.org/documentation/biom_conversion.html>

format.org/documentation/biom_conversion.html.

--centroids *filename*

Output cluster centroid sequences to *filename*, in fasta format. The centroid is the sequence that seeded the cluster (i.e. the first sequence of the cluster).

--clusterout_id

Add cluster identifier information to the output files when using the --centroids, --consout and --profile options.

--clusterout_sort

Sort some output files by decreasing abundance instead of input order. It applies to the --consout, --msaout, --profile, --centroids, and --uc options. For --uc, the sorting applies only to the centroid information part (the C lines).

--cluster_fast *filename*

Clusterize the fasta sequences in *filename*, automatically sort by decreasing sequence length beforehand.

--cluster_size *filename*

Clusterize the fasta sequences in *filename*, automatically sort by decreasing sequence abundance beforehand.

--cluster_smallmem *filename*

Clusterize the fasta sequences in *filename* without automatically modifying their order beforehand. Sequence are expected to be sorted by decreasing sequence length, unless --usersort is used.

--cluster_unoise *filename*

Perform denoising of the fasta sequences in *filename* according to the UNOISE version 3 algorithm by Robert Edgar, but without the *de novo* chimera removal step, which may be performed afterwards with --uchime3_denovo. The options --minsize (default 8) and --unoise_alpha (default 2.0) may be specified. In the this algorithm, clustering of sequences depend on both the sequence distance and the abundance ratio. The abundance ratio (skew) is the abundance of a new sequence divided by the abundance of the centroid sequence. This skew must not be larger than beta if the sequences should be clustered together. Beta is calculated as 2 raised to the power of minus 1 minus alpha times the sequence distance. The sequence distance used is the number of mismatches in the alignment, ignoring gaps. This means that the abundance must be exponentially lower as the distance increases from the centroid for a new sequence to be included in the cluster. Nearer sequences with higher abundances will form their own new clusters.

--clusters *string*

Output each cluster to a separate fasta file using the prefix *string* and a ticker (0, 1, 2, etc.) to construct the path and filenames.

--consout *filename*

Output cluster consensus sequences to *filename*. For each cluster, a center-star multiple sequence alignment is computed with the centroid as the center, using a fast algorithm (not accurate when using low pairwise identity thresholds). A consensus sequence is constructed by taking the majority symbol (nucleotide or gap) from each column of the alignment. Columns containing a majority of gaps are skipped, except for terminal gaps. If the --sizein option is specified, sequence abundances will be taken into account.

--cons_truncate

This command is ignored. A warning is issued.

--gapext *string*

Set penalties for a gap extension. See --gapopen for a complete description of the penalty declaration system. The default is to initialize the six gap extending penalties

using a penalty of 2 for extending internal gaps and a penalty of 1 for extending terminal gaps, in both query and target sequences (i.e. 2I/1E).

--gapopen *string*

Set penalties for a gap opening. A gap opening can occur in six different contexts: in the query (Q) or in the target (T) sequence, at the left (L) or right (R) extremity of the sequence, or inside the sequence (I). Sequence symbols (Q and T) can be combined with location symbols (L, I, and R), and numerical values to declare penalties for all possible contexts: aQL/bQI/cQR/dTL/eTI/fTR, where abcdef are zero or positive integers, and '/' is used as a separator.

To simplify declarations, the location symbols (L, I, and R) can be combined, the symbol (E) can be used to treat both extremities (L and R) equally, and the symbols Q and T can be omitted to treat query and target sequences equally. For instance, the default is to declare a penalty of 20 for opening internal gaps and a penalty of 2 for opening terminal gaps (left or right), in both query and target sequences (i.e. 20I/2E). If only a numerical value is given, without any sequence or location symbol, then the penalty applies to all gap openings. To forbid gap-opening, an infinite penalty value can be declared with the symbol '**'. To use **vsearch** as a semi-global aligner, a null-penalty can be applied to the left (L) or right (R) gaps.

vsearch always initializes the six gap opening penalties using the default parameters (20I/2E). The user is then free to declare only the values he/she wants to modify. The *string* is scanned from left to right, accepted symbols are (0123456789/LIREQT*), and later values override previous values.

Please note that **vsearch**, in contrast to **usearch**, only allows integer gap penalties. Because the lowest gap penalties are 0.5 by default in **usearch**, all default scores and gap penalties in **vsearch** have been doubled to maintain equivalent penalties and to produce identical alignments.

--id *real* Do not add the target to the cluster if the pairwise identity with the centroid is lower than *real* (value ranging from 0.0 to 1.0 included). The pairwise identity is defined as the number of (matching columns) / (alignment length - terminal gaps). That definition can be modified by --iddef.

--iddef 0|1|2|3|4

Change the pairwise identity definition used in --id. Values accepted are:

0. CD-HIT definition: (matching columns) / (shortest sequence length).
1. edit distance: (matching columns) / (alignment length).
2. edit distance excluding terminal gaps (same as --id).
3. Marine Biological Lab definition counting each gap opening (internal or terminal) as a single mismatch, whether or not the gap was extended: 1.0 - [(mismatches + gap openings)/(longest sequence length)]
4. BLAST definition, equivalent to --iddef 1 in a context of global pairwise alignment.

--lengthout

Write sequence length information to the output files in FASTA format by adding a ";length=*integer*" attribute in the header.

--match *integer*

Score assigned to a match (i.e. identical nucleotides) in the pairwise alignment. The default value is 2.

--minsize *positive integer*

Specify the minimum abundance of sequences for denoising using --cluster_unequal. The default is 8.

--mismatch *integer*

Score assigned to a mismatch (i.e. different nucleotides) in the pairwise alignment. The default value is -4.

--msaout *filename*

Output a multiple sequence alignment and a consensus sequence for each cluster to *filename*, in fasta format. Be warned that vsearch computes center star multiple sequence alignments using a fast method whose accuracy can decrease significantly when using low pairwise identity thresholds. The consensus sequence is constructed by taking the majority symbol (nucleotide or gap) from each column of the alignment. Columns containing a majority of gaps are skipped, except for terminal gaps. If the --sizein option is specified, sequence abundances will be taken into account when computing the consensus.

--mothur_shared_out *filename*

Output an OTU table in the mothur 'shared' tab-separated plain text format as described at ([link](#)) <https://www.mothur.org/wiki/Shared_file>. The format describes how a matrix containing the abundances of the OTUs in the different samples is stored. The first line will start with the strings 'label', 'group' and 'numOtus' and is followed by a list of all OTU identifiers. The following lines, one for each sample, starts with the string 'vsearch' followed by the sample identifier, the total number of OTUs, and a list of abundances for each OTU in that sample, in the order given on the first line. The OTU and sample identifiers are extracted from the FASTA headers of the sequences. The OTUs are represented by the cluster centroids. See the --biomout option for further details.

--otutabout *filename*

Output an OTU table in the classic tab-separated plain text format as a matrix containing the abundances of the OTUs in the different samples. The first line will start with the string '#OTU ID' and is followed by a tab-separated list of all sample identifiers. The following lines, one for each OTU, starts with the OTU identifier and is followed by a tab-separated list of abundances for that OTU in each sample, in the order given on the first line. The OTU and sample identifiers are extracted from the FASTA headers of the sequences (see the --sample option). The OTUs are represented by the cluster centroids. An extra column is added to the right of the table if taxonomy information is available for at least one of the OTUs. This column will be labelled 'taxonomy' and each row will then contain the taxonomy information extracted for that OTU. See the --biomout option for further details.

--profile *filename*

Output a sequence profile to a text file with the frequency of each nucleotide in each position in the multiple alignment for each cluster. There is a FASTA-like header line for each cluster, followed by the profile information in a tab-separated format. The eight columns are: position (0-based), consensus nucleotide, number of As, number of Cs, number of Gs, number of Ts or Us, number of gap symbols, and finally the total number of ambiguous nucleotide symbols (B, D, H, K, M, N, R, S, Y, V or W). All numbers are integers. If the --sizein option is specified, sequence abundances will be taken into account.

--qmask *none|dust|soft*

Mask regions in sequences using the *dust* or the *soft* methods, or do not mask (*none*). Warning, when using *soft* masking, clustering becomes case sensitive. The default is to mask using *dust*.

--qsegout *filename*

Write the aligned part of each query sequence to *filename* in FASTA format.

--relabel *string*

Relabel sequence identifiers in the output files produced by --consout, --profile and --centroids options. Please see the description of the same option under Chimera detection for details.

--relabel_keep

When relabelling, keep the old identifier in the header after a space.

--relabel_md5

Relabel sequence identifiers in the output files produced by --consout, --profile and --centroids options. Please see the description of the same option under Chimera detection for details.

--relabel_self

Relabel sequence identifiers in the output files produced by --consout, --profile and --centroids options. Please see the description of the same option under Chimera detection for details.

--relabel_sha1

Relabel sequence identifiers in the output files produced by --consout, --profile and --centroids options. Please see the description of the same option under Chimera detection for details.

--sizein Take into account the abundance annotations present in the input fasta file (search for the pattern '[>]size=*integer*[;]' in sequence headers).

--sizeorder

When an amplicon is close to 2 or more centroids, both within the distance specified with the --id option, resolve the ambiguity by clustering it with the centroid having the highest abundance, not necessarily the closest one. The option only has effect when the value specified with --maxaccepts is higher than one. The --sizeorder option turns on what is sometimes referred to as abundance-based greedy clustering (AGC), in contrast to the default distance-based greedy clustering (DGC).

--sizeout Add abundance annotations to the output fasta files (add the pattern ';size=*integer*;' to sequence headers). If --sizein is specified, abundance annotations are reported to output files, and each cluster centroid receives a new abundance value corresponding to the total abundance of the amplicons included in the cluster (--centroids option). If --sizein is not specified, input abundances are set to 1 for amplicons, and to the number of amplicons per cluster for centroids.

--strand *plus|both*

When comparing sequences with the cluster seed, check the *plus* strand only (default) or check *both* strands.

--tsegout *filename*

Write the aligned part of each target sequence to *filename* in FASTA format.

--uc *filename*

Output clustering results in *filename* using a tab-separated uclust-like format with 10 columns and 3 different type of entries (S, H or C). Each fasta sequence in the input file can be either a cluster centroid (S) or a hit (H) assigned to a cluster. Cluster records (C) summarize information (size, centroid label) for each cluster. In the context of clustering, the option --uc_allhits has no effect on the --uc output. Column content varies with the type of entry (S, H or C):

1. Record type: S, H, or C.
2. Cluster number (zero-based).

3. Centroid length (S), query length (H), or cluster size (C).
4. Percentage of similarity with the centroid sequence (H), or set to '*' (S, C).
5. Match orientation + or - (H), or set to '*' (S, C).
6. Not used, always set to '*' (S, C) or to zero (H).
7. Not used, always set to '*' (S, C) or to zero (H).
8. set to '*' (S, C) or, for H, compact representation of the pairwise alignment using the CIGAR format (Compact Idiosyncratic Gapped Alignment Report): M (match/mismatch), D (deletion) and I (insertion). The equal sign '=' indicates that the query is identical to the centroid sequence (ignoring terminal gaps).
9. Label of the query sequence (H), or of the centroid sequence (S, C).
10. Label of the centroid sequence (H), or set to '*' (S, C).

--unoise_alpha real

Specify the alpha parameter to the --cluster_unequal command. The default is 2.0.

--usersort

When using --cluster_smallmem, allow any sequence input order, not just a decreasing length ordering.

--xlength Strip header attribute ";length=integer" from input sequences. This attribute is added to output sequences by the --lengthout option.**--xsize** Strip abundance information from the headers when writing the output file.

... Most searching options as well as score filtering, gap penalties and masking also apply to clustering (see the Searching section for definitions): --alnout, --blast6out, --fastapairs, --matched, --notmatched, --maxaccepts, --maxrejects, --samout, --userout, --userfields

Dereplication and rereplication options:

VSEARCH can derePLICATE sequences with the commands --derep_fulllength, --derep_id, --derep_smallmem, --derep_prefix and --fastx_uniques. The --derep_fulllength command is deprecated and is replaced by the new --fastx_uniques command that can also handle FASTQ files in addition to FASTA files. The --derep_fulllength, --derep_smallmem, and --fastx_uniques commands require strictly identical sequences of the same length, but ignores upper/lower case and treats T and U as identical symbols. The --derep_id command requires both identical sequences and identical headers/labels. The --derep_prefix command will group sequences with a common prefix and does not require them to be equally long. The --derep_smallmem uses a much smaller amount of memory when dereplicating than the other files, and may be a bit slower and cannot read the input from a pipe. It takes both FASTA and FASTQ files as input but only writes FASTA output to the file specified with the --fastaout option. The --fastx_uniques command can write FASTQ output (specified with --fastqout) or FASTA output (specified with --fastaout) as well as a special tab-separated column text format (with --tabbedout). The other commands can write FASTA output to the file specified with the --output option. All dereplication commands, except --derep_smallmem, can write output to a special UCLUST-like file specified with the --uc option. The --rereplicate command can duplicate sequences in the input file according to the abundance of each input sequence. Other valid options are --fastq_ascii, --fastq_asciiout, --fastq_qmax, --fastq_qmaxout, --fastq_qmin, --fastq_qminout, --fastq_qout_max, --lengthout, --maxuniquesize, --minuniquesize, --relabel, --relabel_keep, --relabel_md5, --relabel_self, --relabel_sha1, --sizein, --sizeout, --strand, --topn, --xlength, and --xsize.

--deref_fulllength *filename*

Merge strictly identical sequences contained in *filename*. Identical sequences are defined as having the same length and the same string of nucleotides (case insensitive, T and U are considered the same). See the options --sizein and --sizeout to take into account and compute abundance values. This command does not support multithreading.

--deref_id *filename*

Merge strictly identical sequences contained in *filename*, as with the --deref_fulllength command, but the sequence labels (identifiers) on the header line need to be identical too.

--deref_smallmem *filename*

Merge strictly identical sequences contained in *filename*, as with the --deref_fulllength command, but using much less memory. The output is written to a FASTA file specified with the --fastaout option. The output is written in the order that the sequences first appear in the input, and not in descending abundance order as with the other dereplication commands. It can read, but not write FASTQ files. This command cannot read from a pipe, it must be a proper file, as it is read twice. Dereplication is performed with a 128 bit hash function and it is not verified that grouped sequences are identical, however the probability that two different sequences are grouped in a dataset of one billion unique sequences is approximately 1e-21. Memory footprint is appr. 24 bytes times the number of unique sequence. Multithreading and the options --topn, --uc, or --tabbedout are not supported.

--deref_prefix *filename*

Merge sequences with identical prefixes contained in *filename*. A short sequence identical to an initial segment (prefix) of another sequence is considered a replicate of the longer sequence. If a sequence is identical to the prefix of two or more longer sequences, it is clustered with the shortest of them. If they are equally long, it is clustered with the most abundant. Remaining ties are solved using sequence headers and sequence input order. Sequence comparisons are case insensitive, and T and U are considered identical. This command does not support multithreading.

--fastaout *filename*

Write the dereplicated sequences to *filename*, in fasta format and sorted by decreasing abundance. Identical sequences receive the header of the first sequence of their group. If --sizeout is used, the number of occurrences (i.e. abundance) of each sequence is indicated at the end of their fasta header using the pattern '`;size=integer;`'. This option is only valid for --fastx_uniques and --deref_smallmem.

--fastqout *filename*

Write the dereplicated sequences to *filename*, in fastq format and sorted by decreasing abundance. Identical sequences receive the header of the first sequence of their group. If --sizeout is used, the number of occurrences (i.e. abundance) of each sequence is indicated at the end of their fastq header using the pattern '`;size=integer;`'. This option is only valid for --fastx_uniques.

--fastq_ascii *positive integer*

Define the ASCII character number used as the basis for the FASTQ quality score. The default is 33, which is used by the Sanger / Illumina 1.8+ FASTQ format (phred+33). The value 64 is used by the Solexa, Illumina 1.3+ and Illumina 1.5+ formats (phred+64). Only 33 and 64 are valid arguments.

--fastq_asciiout *positive integer*

When using --fastq_convert, --sff_convert or --fasta2fastq, define the ASCII character number used as the basis for the FASTQ quality score when writing FASTQ output files. The default is 33. Only 33 and 64 are valid arguments.

--fastq_qmax *positive integer*

Specify the maximum quality score accepted when reading FASTQ files. The default is 41, which is usual for recent Sanger/Illumina 1.8+ files.

--fastq_qmaxout *positive integer*

Specify the maximum quality score used when writing FASTQ files. The default is 41, which is usual for recent Sanger/Illumina 1.8+ files. Older formats may use a maximum quality score of 40.

--fastq_qmin *positive integer*

Specify the minimum quality score accepted for FASTQ files. The default is 0, which is usual for recent Sanger/Illumina 1.8+ files. Older formats may use scores between -5 and 2.

--fastq_qminout *positive integer*

Specify the minimum quality score used when writing FASTQ files. The default is 0, which is usual for Sanger/Illumina 1.8+ files. Older versions of the format may use scores between -5 and 2.

--fastq_qout_max

For --fastx_uniques, indicate that the new quality scores computed when dereplicating FASTQ files should be equal to the maximum (best) of the input quality scores for each position (corresponding to the lowest error probability). The default is to output a quality score corresponding to the average of the error probabilities for each position.

--fastx_uniques *filename*

Merge strictly identical sequences contained in FASTA or FASTQ file *filename*. Identical sequences are defined as having the same length and the same string of nucleotides (case insensitive, T and U are considered the same). See the options --sizein and --sizeout to take into account and compute abundance values. This command does not support multithreading. By default, the quality scores in FASTQ output files will correspond to the average error probability of the nucleotides in the each position. If the --fastq_qout_max option is given, the quality score will be the highest (best) quality score observed in each position.

--lengthout

Write sequence length information to the output files in FASTA and FASTQ format by adding a ";length=*integer*" attribute in the header.

--maxuniquesize *positive integer*

Discard sequences with a post-dereplication abundance value greater than *integer*.

--minuniquesize *positive integer*

Discard sequences with a post-dereplication abundance value smaller than *integer*.

--output *filename*

Write the dereplicated sequences to *filename*, in fasta format and sorted by decreasing abundance. Identical sequences receive the header of the first sequence of their group. If --sizeout is used, the number of occurrences (i.e. abundance) of each sequence is indicated at the end of their fasta header using the pattern ';size=*integer*;'. This option is not allowed for --fastx_uniques or --deref_smallmem.

--relabel *string*

Please see the description of the same option under Chimera detection for details.

--relabel_keep

When relabelling, keep the old identifier in the header after a space.

--relabel_md5

Please see the description of the same option under Chimera detection for details.

--relabel_self

Please see the description of the same option under Chimera detection for details.

--relabel_sha1

Please see the description of the same option under Chimera detection for details.

--rereplicate *filename*

Duplicate each sequence the number of times indicated by the abundance of each sequence in the specified file (option --sizein is always implied). The sequence labels are identical for the same sequence, unless --relabel, --relabel_self, --relabel_sha1 or --relabel_md5 is used to create unique labels. Output is written to the file specified with the --output option, in FASTA format. The output file does not contain abundance information unless --sizeout is specified, in which case an abundance of 1 is used.

--sizein Take into account the abundance annotations present in the input fasta file (search for the pattern '[>]size=*integer*[;]' in sequence headers). That option is active by default when rereplicating.

--sizeout Add abundance annotations to the output fasta file (add the pattern ';size=*integer*;' to sequence headers). If --sizein is specified, each unique sequence receives a new abundance value corresponding to its total abundance (sum of the abundances of its occurrences). If --sizein is not specified, input abundances are set to 1, and each unique sequence receives a new abundance value corresponding to its number of occurrences in the input file.

--strand *plus|both*

When searching for strictly identical sequences, check the *plus* strand only (default) or check *both* strands.

--tabbedout *filename*

Output clustering info to the specified tab-separated text file with 6 columns and a row for each input sequence. Column 1 contains the original label/header of the sequence. Column 2 contains the label of the output sequence which is equal to the label/header of the first sequence in each cluster, but potentially relabelled. Column 3 contains the cluster number, starting from 0. Column 4 contains the sequence number within each cluster, starting at 0. Column 5 contains the number of sequences in the cluster. Column 6 contains the original label/header of the first sequence in the cluster before any potential relabelling. This option is only valid for the --fastx_uniques command.

--topn *positive integer*

Output only the top *integer* sequences (i.e. the most abundant).

--uc *filename*

Output full-length or prefix-dereplication results in *filename* using a tab-separated uclust-like format with 10 columns and 3 different type of entries (S, H or C). Each fasta sequence in the input file can be either a cluster centroid (S) or a hit (H) assigned to a cluster. Cluster records (C) summarize information (size, centroid label) for each cluster. In the context of dereplication, the option --uc_allhits has no effect on the --uc output. Column content varies with the type of entry (S, H or C):

1. Record type: S, H, or C.
2. Cluster number (zero-based).
3. Sequence length (S, H), or cluster size (C).
4. Percentage of similarity with the centroid sequence (H), or set to '*' (S, C).
5. Match orientation + or - (H), or set to '*' (S, C).

6. Not used, always set to '*' (S, C) or 0 (H).
7. Not used, always set to '*' (S, C) or 0 (H).
8. Not used, always set to '*'.
9. Label of the query sequence (H), or of the centroid sequence (S, C).
10. Label of the centroid sequence (H), or set to '*' (S, C).

--xlength

Strip header attribute ";length=*integer*" from input sequences. This attribute is added to output sequences by the --lengthout option.

--xsize Strip abundance information from the headers when writing the output file.

Extraction options:

Sequences with headers matching certain criteria can be extracted from FASTA and FASTQ files using the --fastx_getseq, --fastx_getseqs and --fastx_getsubseq commands.

The --fastx_getseq command requires the header to match a label specified with the --label option. If the --label_substr_match option is given, the label may be a substring located anywhere in the header, otherwise the entire header must match the label. These matches are not case-sensitive. The headers in the input file are truncated at the first space or tab character unless the --notrunclabels option is given. The matching sequences will be written to the files specified with the --fastaout and --fastqout options, in FASTA and FASTQ format, respectively. Sequences that do not match are written to the files specified with the --notmatched and --notmatchedfq options, respectively.

The --fastx_getsubseq command is similar to the --fastx_getseq command, but will extract a subsequence of the matching sequences. The start position is specified with the --subseq_start option and the end position is specified with the --subseq_end option. The positions are 1-based, meaning that the first symbol of the sequence is at position 1. If the start or end position option is not specified, the default is to start at the first position and end at the last position in the sequence.

The --fastx_getseqs command is similar to the --fastx_getseq command but allows more flexibility in specifying the label(s) to be matched. A single label may be specified using the --label option as described above. Alternatively, a file containing a list of labels to be matched may be specified with the --labels option. The file must be a plain text file with one label on each line. The --label_word and --label_words options may be used to specify either a single word or a file containing a list of words, respectively, to be matched. Words are defined as character sequences delimited either by a character that is not alpha-numeric (A-Z, a-z, or 0-9) or by the beginning or end of the header. Word matching is case-sensitive. The --label_field option will limit the matching of words to a certain field in the header.

--fastaout *filename*

Write the extracted sequences in FASTA format to the file with the given name.

--fastqout *filename*

Write the extracted sequences in FASTQ format to the file with the given name. This option is illegal if the input is in FASTA format.

--fastx_getseq *filename*

Extract sequences from the given FASTA or FASTQ file. Specify a label to match using the --label option. Output files are specified with the --fastaout, --fastqout, --notmatched and --notmatchedfq options.

--fastx_getseqs *filename*

Extract sequences from the given FASTA or FASTQ file. Specify the label or labels to match using one of the following options: --label, --labels, --label_word, or --label_words. Output files are specified with the --fastaout, --fastqout, --notmatched and --notmatchedfq options.

--fastx_getsubseq *filename*

Extract a certain part of some of the sequences in the given FASTA or FASTQ file. Specify labels to match using the --label option. Specify the subsequence range to be extracted with the --subseq_start and --subseq_end options. Output files are specified with the --fastaout, --fastqout, --notmatched and --notmatchedfq options.

--label *string*

Specify the label to match in the sequence header. Unless the --label_substr_match option is given, the label must match the entire header. The comparison is not case-sensitive.

--label_field *string*

Specify a field name to be used when matching using the --label_word or --label_words option. The field name is a string like "abc" that must precede the word to be matched with an equals sign (=) in between. The field must be delimited by semicolons or the beginning or end of the header. The following header will match the label 123 in the field abc: "seq1;abc=123".

--label_substr_match

The labels specified with the --label or the --labels option may match anywhere in the header if this option is given. Otherwise a label needs to match the entire header.

--label_word *string*

Specify a word to match in the sequence header. Words are defined as strings delimited by either the start or end of the header or by any symbol that is not a letter (A-Z, a-z) or digit (0-9). The comparison is case-sensitive.

--label_words *filename*

Specify a file containing words to be matched against the sequence headers. The plain text file must contain one word on each line. Words are defined as strings delimited by either the start or end of the header or by any symbol that is not a letter (A-Z, a-z) or digit (0-9). The comparison is case-sensitive.

--labels *filename*

Specify a file containing labels to be matched against the sequence headers. The plain text file must contain one label on each line. Unless the --label_substr_match option is given, a label must match the entire header. The comparison is not case-sensitive.

--notmatched *filename*

Write the sequences that were not extracted to the file with the given name, in FASTA format.

--notmatchedfq *filename*

Write the sequences that were not extracted to the file with the given name, in FASTQ format. This option is illegal if the input is in FASTA format.

--subseq_end *positive integer*

Specify the end position in the sequences when extracting subsequences using the --fastx_getsubseq command. Positions are 1-based, so the sequences start at position 1. The default is to end at the end of the sequence if this option is not specified.

--subseq_start *positive integer*

Specify the starting position in the sequences when extracting subsequences using the --fastx_getsubseq command. Positions are 1-based, so the sequences start at position 1. The default is to start at the beginning of the sequence (position 1), if this option is not specified.

FASTA/FASTQ/SFF file processing options:

Analyse, trim, filter, convert, merge, join or reverse complement sequences in FASTA, FASTQ or SFF files. The --fastq_chars command can be used to analyse FASTQ files to identify the quality encoding and the range of quality score values used. To convert between different FASTQ file

variants, use the --fastq_convert command. Statistical analysis of the quality and length of the sequences in a FASTQ file may be performed with the --fastq_stats, --fastq_eestats, and --fastq_eestats2 commands. Sequences may be trimmed, filtered and converted by the --fastq_filter or --fastx_filter commands. The --sff_convert command can be used to convert SFF files to FASTQ, while the --fasta2fastq command will convert a FASTA file to a FASTQ file with fake quality scores. Paired-end reads can be merged using the --fastq_mergepairs command or joined with the --fastq_join command. The --fastx_revcomp command will reverse-complements sequences.

--eeout When using --fastq_filter, --fastx_filter or --fastq_mergepairs, include the number of expected errors (ee) in the sequence header of FASTQ and FASTA output files. This option is a synonym of the --fastq_eeout option. Use the --xee option to remove this information from headers.

--eatabbedout *filename*

When specified with the --fastq_mergepairs command, write statistics with expected errors of each merged read to the given file. The file is a tab separated file with four columns: The number of expected errors in the forward read, the number of expected errors in the reverse read, the number of observed errors in the forward read, and the number of observed errors in the reverse read. The observed number of errors are the number of differences in the overlap region of the merged sequence relative to each of the reads in the pair.

--fasta2fastq *filename*

Add a fake nucleotide quality score to the sequences in the given FASTA file and write them to the FASTQ file specified with the --fastqout option. The quality score may be adjusted using the --fastq_qmaxout option (default 41). The --fastq_asciiout option may be used to adjust the FASTQ output quality ASCII base character (default 33).

--fastaout *filename*

When using --fastq_filter, --fastq_mergepairs or --fastx_filter, write to the given FASTA-formatted file the sequences passing the filter, or the merged sequences.

--fastaout_rev *filename*

When using --fastq_filter, or --fastx_filter, write to the given FASTA-formatted file the reverse reads passing the filter.

--fastaout_notmerged_fwd *filename*

When using --fastq_mergepairs, write forward reads not merged to the specified FASTA file.

--fastaout_notmerged_rev *filename*

When using --fastq_mergepairs, write reverse reads not merged to the specified FASTA file.

--fastaout_discarded *filename*

Write sequences that do not pass the filter of the --fastq_filter or --fastx_filter command to the given FASTA-formatted file.

--fastaout_discarded_rev *filename*

Write reverse reads that do not pass the filter of the --fastq_filter or --fastx_filter command to the given FASTA-formatted file.

--fastq_allowmergestagger

When using --fastq_mergepairs, allow merging of staggered read pairs. Staggered pairs are pairs where the 3' end of the reverse read has an overhang to the left of the 5' end of the forward read. This situation can occur when a very short fragment is sequenced. The 3' overhang of the reverse read is not included in the merged sequence. The opposite option is the --fastq_nostagger option. The default is to discard staggered pairs.

--fastq_ascii *positive integer*

Define the ASCII character number used as the basis for the FASTQ quality score. The default is 33, which is used by the Sanger / Illumina 1.8+ FASTQ format (phred+33). The value 64 is used by the Solexa, Illumina 1.3+ and Illumina 1.5+ formats (phred+64). Only 33 and 64 are valid arguments.

--fastq_asciiout *positive integer*

When using --fastq_convert, --sff_convert or --fasta2fastq, define the ASCII character number used as the basis for the FASTQ quality score when writing FASTQ output files. The default is 33. Only 33 and 64 are valid arguments.

--fastq_chars *filename*

Summarize the composition of sequence and quality strings contained in the input FASTQ file. For each sequence symbol, --fastq_chars gives the number of occurrences of the symbol, its relative frequency and the length of the longest run of that symbol. For each character present in the quality strings, --fastq_chars gives the ASCII value of the character, its relative frequency, and the number of times a *k*-mer of that character appears at the end of quality strings. The length of the *k*-mer can be set using --fastq_tail (4 by default). The command --fastq_chars tries to automatically detect the quality encoding (Solexa, Illumina 1.3+, Illumina 1.5+ or Illumina 1.8+/Sanger) by analyzing the range of observed quality score values. In case of success, --fastq_chars suggests values for the --fastq_ascii (33 or 64), --fastq_qmin and --fastq_qmax options to be used with the other commands that require a FASTQ input file.

--fastq_convert *filename*

Convert between the different variants of the FASTQ file format. The quality encoding of the input file must be specified with the --fastq_ascii option (either 33 or 64, the default is 33), and the output quality encoding must be specified with the --fastq_asciiout option (default 33). The minimum and maximum output quality scores may be limited using the --fastq_qminout and --fastq_qmaxout options. The output file is specified with the --fastqout option.

--fastq_eeout

When using --fastq_filter, --fastx_filter or --fastq_mergepairs, include the number of expected errors (ee) in the sequence header of FASTQ and FASTA files. This option is a synonym of the --eeout option. Use the --xee option to remove this information from headers.

--fastq_eestats *filename*

Analyze a FASTQ file and report statistics on the distributions of quality scores, error probabilities and expected accumulated errors. The report, a table of 21 tab-separated columns, is written to the file specified with the --output option. The first column corresponds to the position in the reads (Pos). The second and third columns correspond to the number of reads (Reads) and percentage of reads (PctRecs) that include this position. The remaining columns include information about the distribution of quality scores in this position (Q), error probabilities in this position (Pe), and finally the expected number of accumulated errors from the beginning of the reads and until the current position (EE). For each of the Q, Pe and EE distributions, the following statistics are included: minimum value (Min), lower quartile (Low), median (Med), mean (Mean), upper quartile (Hi), and maximum value (Max). The quality encoding and the range of quality values may be specified with --fastq_ascii --fastq_qmin and --fastq_qmax.

--fastq_eestats2 *filename*

Analyze the specified FASTQ file and report statistics on the number of sequences that would be retained at a combination of selected cutoffs for length truncation and maximum expected errors, that could potentially be used as arguments to the --fastq_trunclen and --fastq_maxee options to the --fastq_filter command. The result, a table of

two or more columns, is written to the file specified with the --output option. There is a line for each length truncation cutoff. The first column on each line contains the selected truncation length, while the following columns contain the number of sequences and, in parenthesis, the percentage of sequences that would be retained at the selected EE levels. The truncation length cutoffs may be specified with the --length_cutoffs option and requires a list of three comma-separated integers indicating the shortest cutoff, the longest cutoff, and the increment between cutoffs. The longest cutoff may be specified with a star (*) which indicates that the limit is equal to the longest sequence in the input file. The default setting is "50,*,50" meaning that truncation lengths of 50, 100, 150 and so on up to the longest sequence length should be used. The maximum expected error (EE) cutoffs may be specified with the --ee_cutoffs option which requires a comma-separated list of floating point numbers as its argument. The default setting is "0.5,1.0,2.0" that indicates that expected error levels of 0.5, 1.0 and 2.0 should be used.

--fastq_filter *filename*

Trim and/or filter sequences in the given FASTQ file. Similar to the --fastx_filter command, but works only on FASTQ files. See --fastx_filter for details.

--fastq_join *filename*

Join paired-end sequence reads into one sequence and add a gap between them using a padding sequence. The sequences are not merged as with the fastq_mergepairs command, but simply joined with a gap. The forward reads are specified as the argument to this option and the reverse reads are specified with the --reverse option. The resulting sequences consist of the forward read, the padding sequence and the reverse complement of the reverse read. The padding sequence is specified with the --join_padgap option and the padding quality is specified with the --join_padgapq option. The default padding sequence string is NNNNNNNNN and the default padding quality string is IIIII-III, corresponding to a base quality score of 40 (a very high quality score with error probability 0.00001). The joined sequences are output to the file(s) specified with the --fastaout or --fastqout options.

--fastq_maxdiffs *positive integer*

When using --fastq_mergepairs, specify the maximum number of non-matching nucleotides allowed in the overlap region. That option has a strong influence on the merging success rate. The default value is 10.

--fastq_maxdiffpct *real*

When using --fastq_mergepairs, specify the maximum percentage of non-matching nucleotides allowed in the overlap region. The default value is 100.0%. There are other more sophisticated rules in the merging algorithm that will discard read pairs with a high fraction of mismatches.

--fastq_maxee *real*

When using --fastq_filter, --fastq_mergepairs or --fastx_filter, discard sequences with an expected error greater than the specified number (value ranging from 0.0 to infinity). For a given sequence, the expected error is the sum of error probabilities for all the positions in the sequence. Since error probabilities can be small but not null, the expected error is always greater than zero, and at most equal to the length of the sequence when all positions in the sequence have an error probability of 1.0.

Using the expected error as the *lambda* parameter in the Poisson distribution, it is possible to compute the probability of observing *k* errors. For instance, a read with an expected error of 1.0 has:

- 36.8% chance of having zero error,
- 36.8% chance of having one error,

- 18.4% chance of having two errors,
- 6.1% chance of having three errors,
- 1.5% chance of having four errors,
- 0.3% chance of having five errors,
- etc.

--fastq_maxee_rate *real*

When using --fastq_filter or --fastx_filter, discard sequences with an average expected error greater than the specified number (value ranging from 0.0 to 1.0 included). For a given sequence, the average expected error is the sum of error probabilities for all the positions in the sequence, divided by the length of the sequence.

--fastq maxlen *positive integer*

When using --fastq_filter, --fastq_mergepairs or --fastx_filter, discard sequences with more than the specified number of bases.

--fastq_maxmergelen *positive integer*

When using --fastq_mergepairs, specify the maximum length of the merged sequence (default is 1,000,000).

--fastq_maxns *positive integer*

When using --fastq_filter, --fastq_mergepairs or --fastx_filter, discard sequences with more than the specified number of N's.

--fastq_mergepairs *filename*

Merge paired-end sequence reads into one sequence. The forward reads are specified as the argument to this option and the reverse reads are specified with the --reverse option. Reads with the same index/position in the forward and reverse files are considered to form a pair, even if their labels are different. Thus, forward and reverse reads **must** appear in the same order and total number in both files. A warning is emitted if the forward and reverse files contain different numbers of reads. The merged sequences are written to the file(s) specified with the --fastaout or --fastqout options. The non-merged reads can be output to the files specified with the --fastaout_notmerged_fwd, --fastaout_notmerged_rev, --fastqout_notmerged_fwd and --fastqout_notmerged_rev options. Statistics may be output to the file specified with the --eetabbedout option. Sequences are truncated as specified with the --fastq_truncqual option to remove low-quality bases in the 3' end. Sequences shorter than specified with --fastq_minlen (after truncation) are discarded (1 by default). Sequences with too many ambiguous bases (N's), as specified with the --fastq_maxns are also discarded (no limit by default). Staggered reads are not merged unless the --fastq_allowmergestagger option is specified. The minimum length of the overlap region between the reads may be specified with the --fastq_minovlen option (at least 5, default 10). The overlap region may not include more mismatches than specified with the --fastq_maxdifs option (10 by default) or a higher percentage of mismatches than specified with the --fastq_maxdiffpct option (100.0% by default), otherwise the read pair is discarded. Additional rules will avoid merging of reads that cannot be aligned reliably and unambiguously. The minimum and maximum length of the merged sequence may be specified with the --fastq_minmergelen and --fastq_maxmergelen options, respectively. The quality value limits for output files may be specified with the --fastq_qminout and --fastq_qmaxout options, but they apply only to the merged region. Other relevant options are: --fastq_ascii, --fastq_maxee, --fastq_nostagger, --fastq_qmax, --fastq_qmin, and --label_suffix.

--fastq_minlen *positive integer*

When using --fastq_filter, --fastq_mergepairs or --fastx_filter, discard input sequences with less than the specified number of bases (default 1).

--fastq_minmrgelen *positive integer*

When using --fastq_mergepairs, specify the minimum length of the merged sequence. The default is 1.

--fastq_minovlen *positive integer*

When using --fastq_mergepairs, specify the minimum overlap between the merged reads. The default is 10. Must be at least 5.

--fastq_minqual *positive integer*

When using --fastq_filter or --fastx_filter, discard reads having any base with a quality score below the given value. The default is 0, which discards none.

--fastq_nostagger

When using --fastq_mergepairs, forbid the merging of staggered read pairs. This is the default behaviour of --fastq_mergepairs. To change that behaviour, see the --fastq_lowmergestagger option.

--fastq_qmax *positive integer*

Specify the maximum quality score accepted when reading FASTQ files. The default is 41, which is usual for recent Sanger/Illumina 1.8+ files.

--fastq_qmaxout *positive integer*

When using --fastq_mergepairs, --fastq_convert, --sff_convert or --fasta2fastq, specify the maximum quality score used when writing FASTQ files. For the --fasta2fastq command, the value specified here is the fake quality score used for the FASTQ output file. The default is 41, which is usual for recent Sanger/Illumina 1.8+ files. Older formats may use a maximum quality score of 40. The limit only applies to the merged region when using --fastq_mergepairs.

--fastq_qmin *positive integer*

Specify the minimum quality score accepted for FASTQ files. The default is 0, which is usual for recent Sanger/Illumina 1.8+ files. Older formats may use scores between -5 and 2.

--fastq_qminout *positive integer*

When using --fastq_mergepairs, --fastq_convert or --sff_convert, specify the minimum quality score used when writing FASTQ files. The default is 0, which is usual for Sanger/Illumina 1.8+ files. Older versions of the format may use scores between -5 and 2. The limit applies only to the merged region when using --fastq_mergepairs.

--fastq_stats *filename*

Analyze a FASTQ file and report the number of reads it contains. The quality encoding and the range of quality values may be specified with --fastq_ascii --fastq_qmin and --fastq_qmax. That command requires the --log option and outputs the following detailed statistics on read length, quality score, length vs. quality distributions, and length / quality filtering:

Read length distribution:

1. L: read length.
2. N: number of reads.
3. Pct: fraction of reads with this length.
4. AccPct: fraction of reads with this length or longer.

Quality score distribution:

1. ASCII: character encoding the quality score.
2. Q: Phred quality score.

3. Pe: probability of error associated with the quality score.
4. N: number of bases with this quality score.
5. Pct: fraction of bases with this quality score.
- 6: AccPct: fraction of bases with this quality score or higher.

Length vs. quality distribution:

1. L: position in reads (starting from position 2).
2. PctRecs: fraction of reads with at least this length.
3. AvgQ: average quality score at this position.
4. P(AvgQ): error probability corresponding to AvgQ.
5. AvgP: average error probability at this position.
- 6: AvgEE: average expected error over all reads up to this position.
- 7: Rate: growth rate of AvgEE between this position and position - 1.
- 8: RatePct: Rate (as explained above) expressed as a percentage.

Effect of expected error and length filtering:

The first column indicates read lengths (*L*). The next four columns indicate the number of reads that would be retained by the `--fastq_filter` command if the reads were truncated at length *L* (option `--fastq_trunclen L`) and filtered to have a maximum expected error of 1.0, 0.5, 0.25 or 0.1 (with the option `--fastq_maxee float`). The last four columns indicate the fraction of reads that would be retained by the `--fastq_filter` command using the same length and maximum expected error parameters.

Effect of minimum quality and length filtering:

The first column indicates read lengths (*Len*). The next four columns indicate the fraction of reads that would be retained by the `--fastq_filter` command if the reads were truncated at length *Len* (option `--fastq_trunclen Len`) or at the first position with a quality *Q* equal to or lesser than 5, 10, 15 or 20 (option `--fastq_truncqual Q`).

--fastq_stripleft positive integer

When using `--fastq_filter` or `--fastx_filter`, strip the specified number of bases from the left end of the reads. If the length of the resulting read is null, then the read is discarded.

--fastq_stripright positive integer

When using `--fastq_filter` or `--fastx_filter`, strip the specified number of bases from the right end of the reads. If the length of the resulting read is null, then the read is discarded.

--fastq_tail positive integer

When using `--fastq_chars`, count the number of times a series of characters of length *k* appears at the end of quality strings. By default, *k* = 4.

--fastq_truncee real

When using `--fastq_filter` or `--fastx_filter`, truncate sequences so that their total expected error is not higher than the specified value.

--fastq_truncee_rate real

When using `--fastq_filter` or `--fastx_filter`, truncate sequences so that their average expected error per base is not higher than the specified value. The truncation will happen at the first occurrence. The average expected error per base is calculated as the total expected number of errors divided by the length of the sequence after truncation.

--fastq_trunclen *positive integer*

When using --fastq_filter or --fastx_filter, truncate sequences to the specified length. Shorter sequences are discarded.

--fastq_trunclen_keep *positive integer*

When using --fastq_filter or --fastx_filter, truncate sequences to the specified length. Shorter sequences are not discarded.

--fastq_truncqual *positive integer*

When using --fastq_filter, --fastq_mergepairs or --fastx_filter, truncate sequences starting from the first base with the specified base quality score value or lower.

--fastqout *filename*

When using --fastq_filter, --fastq_mergepairs, --fastx_filter or --fasta2fastq, write to the given FASTQ-formatted file the sequences passing the filter, or the merged or converted sequences.

--fastqout_rev *filename*

When using --fastq_filter or --fastx_filter, write to the given FASTQ-formatted file the reverse reads passing the filter.

--fastqout_discarded *filename*

When using --fastq_filter or --fastx_filter, write sequences that do not pass the filter to the given FASTQ-formatted file.

--fastqout_discarded_rev *filename*

When using --fastq_filter or --fastx_filter, write reverse reads that do not pass the filter to the given FASTQ-formatted file.

--fastqout_notmerged_fwd *filename*

When using --fastq_mergepairs, write forward reads not merged to the specified FASTQ file.

--fastqout_notmerged_rev *filename*

When using --fastq_mergepairs, write reverse reads not merged to the specified FASTQ file.

--fastx_filter *filename*

Trim and/or filter the sequences in the given FASTA or FASTQ file and output the remaining sequences to the FASTQ file specified with the --fastqout option and/or to the FASTA file specified with the --fastaout option. Discarded sequences are written to the files specified with the --fastaout_discarded and --fastqout_discarded options. The input format (FASTA or FASTQ) is automatically detected. If the input consists of paired sequences, an input file with reverse reads may be specified with the --reverse option, and corresponding output will be written to the files specified with the --fastqout_rev, --fastaout_rev, --fastqout_discarded_rev, and --fastaout_discarded_rev options. Output can not be written to FASTQ files if the input is in FASTA format. The sequences are first trimmed and then filtered based on the remaining bases. Sequences may be trimmed using the options --fastq_stripleft, --fastq_stripright, --fastq_truncee, --fastq_truncee_rate, --fastq_trunclen, --fastq_trunclen_keep and --fastq_truncqual. The sequences may be filtered using the options --fastq_maxee, --fastq_maxee_rate, --fastq_maxlen, --fastq_maxns, --fastq_minlen (default 1), --fastq_minqual, --fastq_trunclen, --maxsize, and --minsize. Sequences not satisfying the requirements are discarded. For pairs of sequences, both sequences in a pair must satisfy the requirements, otherwise both are discarded. If no shortening or filtering options are given, all sequences are written to the output files, possibly after conversion from FASTQ to FASTA format. The --relabel option may be used to relabel the output sequences. The --eeout option may be used to output the expected number of errors in each sequence. After all sequences have been processed, the number of kept and discarded sequences will be shown, as well as how many of the kept sequences were trimmed. When the input is in FASTA format, the following options

are not accepted because quality scores are not available: --eeout, --fastq_ascii, --fastq_eeout, --fastq_maxee, --fastq_maxee_rate, --fastq_minqual, --fastq_out, --fastq_qmax, --fastq_qmin, --fastq_truncee, --fastq_truncee_rate, --fastq_truncqual, --fastqout_discarded, --fastqout_discarded_rev, --fastqout_rev.

--fastx_revcomp *filename*

Reverse-complement the sequences in the given FASTA or FASTQ file to a file specified with the --fastaout and/or --fastqout options. If the input file is in FASTA format, the output can not be written back to a FASTQ file due to missing base quality scores.

--join_padgap *string*

When running --fastq_join, use the *string* as a sequence padding string. The default is NNNNNNNN (8 N's). Option accepts an empty string, or any other combination of ASCII characters.

--join_padgapq *string*

When running --fastq_join, use the *string* as a quality padding string. The default is a string of I's equal in length to the sequence padding string. The letter I corresponds to a base quality score of 40 indicating a very high quality base with error probability of 0.0001. Option accepts an empty string, or any other combination of ASCII characters.

--lengthout

Write sequence length information to the output files in FASTA or FASTQ format by adding a ";length=*integer*" attribute in the header.

--maxsize *positive integer*

When using --fastq_filter or --fastx_filter, discard sequences with an abundance higher than the specified value.

--minsize *positive integer*

When using --fastq_filter or --fastx_filter, discard sequences with an abundance lower than the specified value.

--output *filename*

When using --fastq_eestats or --fastq_eestats2, write tabulated results to *filename*. See --fastq_eestats's and --fastq_eestats2's documentation for a complete description of the table.

--relabel_keep

When using --relabel, keep the old identifier in the header after a space.

--relabel *string*

Please see the description of the same option under Chimera detection for details.

--relabel_md5

Please see the description of the same option under Chimera detection for details.

--relabel_self

Please see the description of the same option under Chimera detection for details.

--relabel_sha1

Please see the description of the same option under Chimera detection for details.

--reverse *filename*

When using --fastq_filter, --fastx_filter, --fastq_mergepairs or --fastq_join, specify the FASTQ file containing containing the reverse reads.

--sff_convert *filename*

Convert the given SFF file to FASTQ. The FASTQ output file is specified with the --fastqout option. The sequence may be clipped as specified in the SFF file if the option --sff_clip is specified, otherwise no clipping occurs. Bases that would have been clipped are converted to lower case, while the rest is in upper case. The output quality encoding may be specified with the --fastq_asciiout option (default 33). The minimum and

maximum output quality scores may be limited using the --fastq_qminout and --fastq_qmaxout options.

--sff_clip

Specifies that the sequences converted by the --sff_convert command should be clipped in both ends as indicated in the SFF file. By default no clipping is performed.

--xlength

Strip header attribute ";length=*integer*" from input sequences. This attribute is added to output sequences by the --lengthout option.

--xsize

Strip abundance information from the headers when writing the output file.

--xee

Strip information about expected errors (ee) from the output file headers. This information is added by the --fastq_eeout and --eeout options.

Masking options:

An input sequence can be composed of lower- or uppercase letters. When soft masking is specified, lower case letters are treated as symbols that should be masked. Otherwise the case of the input sequences is ignored.

DUST masking is performed automatically by the commands for chimera detection (uchime_denovo, uchime_ref), clustering (cluster_fast, cluster_smallmem, cluster_size), masking (maskfasta, fastx_mask), pairwise alignment (allpairs_global) and searching (search_exact, usearch_global).

Masking is usually specified with the --qmask option, while the --dbmask option is used for the database sequences specified with the --db option with the --usearch_global, --search_exact and --uchime_ref commands.

The argument to the --qmask and --dbmask option may be none, soft or dust. If the argument is none, no masking is performed. If the argument is soft the lower case symbols are masked. Finally, if the argument is dust, the sequence is masked using the DUST algorithm by Tatusov and Lipman to mask low-complexity regions.

If the --hardmask option is specified, all masked regions are converted to N's, otherwise masked regions are indicated by lower case letters.

If any sequence is masked, the masked version of the sequence (with lower case letters or N's) is used in all output files. Otherwise the sequence is unmodified. The exception is the sequences in the output file specified with the --uchimealns option, where the input sequences are converted to upper case first and lower case letters indicate disagreement between the aligned sequences.

The --qmask option (or --dbmask for database sequences) may be combined with the --hardmask option. The results of using the none, dust or soft argument to --qmask or --dbmask are presented below, assuming each input sequence contains both lower and uppercase symbols.

Results if the --hardmask option is off (default):

- none:** no masking, all symbols used, no change
- dust:** masked symbols lowercased, rest uppercased
- soft:** lowercase symbols masked, no case changes

Results if the --hardmask option is on:

- none:** no masking, all symbols used, no change
- dust:** masked symbols changed to Ns, rest unchanged
- soft:** lowercase symbols masked and changed to Ns

When a sequence region is masked, words in the region are not included in the indices used in the heuristic search algorithm. In all other aspects, the region is treated as other regions.

Regions in sequences that are hardmasked (with N's) have a zero alignment score and do not contribute to an alignment.

--fastaout *filename*

Write the masked sequences to *filename*, in fasta format. Applies only to the --fastx_mask command.

--fastqout *filename*

Write the masked sequences to *filename*, in fastq format. Applies only to the --fastx_mask command.

--fastx_mask *filename*

Mask regions in sequences contained in the specified fasta or fastq file. The default is to mask using DUST (use --qmask to modify that behaviour). The output files are specified with the --fastaout and --fastqout options. The minimum and maximum percentage of unmasked residues may be specified with the --min_unmasked_pct and --max_unmasked_pct options, respectively.

--hardmask

Symbols in masked regions are replaced by N's. The default is to replace the masked regions by lower case letters.

--maskfasta *filename*

Mask regions in sequences contained in the fasta file *filename*. The default is to mask using *dust* (use --qmask to modify that behaviour). The output file is specified with the --output option. This command is deprecated, please use --fastx_mask instead.

--max_unmasked_pct *real*

Discard sequences with more than the specified maximum percentage of unmasked residues. Works only with --fastx_mask.

--min_unmasked_pct *real*

Discard sequences with less than the specified minimum percentage of unmasked residues. Works only with --fastx_mask.

--output *filename*

Write the masked sequences to *filename*, in fasta format. Applies only to the --mask_fasta command.

--qmask *none|dust|soft*

If the argument is dust, mask regions in sequences using the *DUST* algorithm that detects simple repeats and low-complexity regions. This is the default for chimera detection, clustering, masking, pairwise alignment, and searching. If the argument is soft, mask the lower case letters in the input sequence. If the argument is none, do not mask.

Orienting options:

The --orient command can be used to orient the sequences in a given file in either the forward or the reverse complementary direction based on a reference database specified with the --db option. The two strands of each input sequence are compared to the reference database using nucleotide words. If one of the strands shares many more words with at least one sequence in the database than the other, that strand is chosen. The correctly oriented sequences may be written to a FASTA file specified with the --fastaout, and to a FASTQ file specified with the --fastqout option (as long as the input was also in FASTQ format). If the result is uncertain, because the number of matching words is too similar, the original sequence is written to the file specified with the --notmatched option. The results may also be written to a tab-delimited text file specified with the --tabbedout option. This file will contain the query label, the direction (+, - or ?), the number of matching words on the forward strand, and the number of matching words on the reverse complementary strand. By default, a word length of 12 is used for this command. The word length may be adjusted using the --wordlength option. There has to be at least 4 times as many matches on one strand than the other for a strand to be selected. In addition to the common options, the following options may also be specified for this command: --dbmask, --qmask, --relabel, --relabel_keep, --relabel_md5, --relabel_self, --relabel_sha1, --sizein, and --sizeout.

--db *filename*

Read the reference database from the given file. It may be in FASTA, FASTQ or UDB format. If an UDB file is used it should have been created with a wordlength of 12.

--fastaout *filename*

Write the correctly oriented sequences to *filename*, in fasta format.

--fastqout *filename*

Write the correctly oriented sequences to *filename*, in fastq format.

--notmatched *filename*

Write the sequences with undetermined direction to *filename*, in the original format.

--orient *filename*

Orient the sequences in the given file.

--tabbedout *filename*

Write the results to a tab-delimited text file with the specified *filename*. This file will contain the query label, the direction (+, - or ?), the number of matching words on the forward strand, and the number of matching words on the reverse complementary strand.

Pairwise alignment options:

The results of the $n * (n-1) / 2$ pairwise alignments are written to the result files specified with --alnout, --blast6out, --fastapairs --matched, --notmatched, --qsegout, --samout, --tsegout, --uc or --userout (see Searching section below). Specify either the --acceptall option to output all pairwise alignments, or specify an identity level with --id to discard weak alignments. Most other accept/reject options (see Searching options below) may also be used. Sequences are aligned on their *plus* strand only. Masking is performed as usual and specified with --qmask and --hardmask.

--acceptall

Write the results of all alignments to output files. This option overrides all other accept/reject options (including --id).

--allpairs_global *filename*

Perform optimal global pairwise alignments of the fasta sequences contained in *filename*. Each sequence is compared to all sequences that come after it in the file, resulting in a total of $n * (n-1) / 2$ pairwise alignments, where n is the total number of sequences. This command is multi-threaded.

--id *real*

Reject the sequence match if the pairwise identity is lower than *real* (value ranging from 0.0 to 1.0 included).

--threads *positive integer*

Number of computation threads to use (1 to 1024). The number of threads should be lesser or equal to the number of available CPU cores. The default is to use all available resources and to launch one thread per logical core.

--uc *filename*

Output pairwise alignment results in *filename* using a tab-separated uclust-like format with 10 columns. Each sequence is compared to all other sequences, and all hits (--acceptall) or only some hits (--id *float*) are reported, with one pairwise comparison per line:

1. Record type, always set to 'H'.
2. Ordinal number of the target sequence (based on input order, starting from zero).
3. Sequence length.

4. Percentage of similarity with the target sequence.
5. Match orientation, always set to '+'.
6. Not used, always set to zero.
7. Not used, always set to zero.
8. Compact representation of the pairwise alignment using the CIGAR format (Compact Idiosyncratic Gapped Alignment Report): M (match/mismatch), D (deletion) and I (insertion). The equal sign '=' indicates that the query is identical to the centroid sequence (ignoring terminal gaps).
9. Label of the query sequence.
10. Label of the target sequence.

Restriction site cutting options:

The input sequences in the file specified with the --cut command are cut into fragments at all restriction sites matching the pattern given with the --cut_pattern option. The fragments on the forward strand are written to the file specified with the --fastaout file and the fragments on the reverse strand are written to the file specified with the --fastaout_rev option. Input sequences that do not match are written to the file specified with the option --fastaout_discarded, and their reverse complement are also written to the file specified with the --fastaout_discarded_rev option. The relabel options (--relabel, --relabel_self, --relabel_keep, --relabel_md5, and --relabel_sha1) may be used to relabel the output sequences.

--cut *filename*

Specify the input file with sequences in FASTA format.

--cut_pattern *string*

Specify the restriction site cutting pattern and positions. The pattern is a string of lower- or uppercase letters specifying the nucleotides that must match, and may include ambiguous nucleotide symbols. The special characters "^" (circumflex) and "_" (underline) are used to indicate the cutting position on the forward and reverse strand, respectively. For example, the pattern "G^AATT_C" is the pattern for the EcoRI restriction site. For such palindromic patterns (identical to its reverse complement) the command will output all possible fragments on both strands. For non-palindromic sites, it may be necessary to run the command also on the reverse complemented input sequences. Exactly one cutting site on each strand must be indicated.

--fastaout *filename*

Specify the output file for the resulting fragments on the forward strand.

--fastaout_rev *filename*

Specify the output file for the resulting fragments on the reverse strand.

--fastaout_discarded *filename*

Specify the output file for the non-matching sequences.

--fastaout_discarded_rev *filename*

Specify the output file for the non-matching sequences, reverse complemented.

Searching options:

--alnout *filename*

Write pairwise global alignments to *filename* using a human-readable format. Use --rowlen to modify alignment length. Output order may vary when using multiple threads.

--biomout *filename*

Write search results to an OTU table in the biom version 1.0 file format. The query file contains the samples, while the database file contains the OTUs. Sample and OTU identifiers are extracted from the header of these sequences. See the --biomout option

in the Clustering section for further details.

--blast6out *filename*

Write search results to *filename* using a blast-like tab-separated format of twelve fields (listed below), with one line per query-target matching (or lack of matching if --output_no_hits is used). Warning, vsearch uses global pairwise alignments, not blast's seed-and-extend algorithm. Therefore, some common blast output values (alignment start and end, evalue, bit score) are reported differently. Output order may vary when using multiple threads. A similar output can be obtain with --userout *filename* and --userfields query+target+id+alnlen+mism+opens+qlo+qhi+tlo+thi+evalue+bits. A complete list and description is available in the section 'Userfields' of this manual.

1. *query*: query label.
2. *target*: target (database sequence) label. The field is set to '*' if there is no alignment.
3. *id*: percentage of identity (real value ranging from 0.0 to 100.0). The percentage identity is defined as $100 * (\text{matching columns}) / (\text{alignment length} - \text{terminal gaps})$. See fields id0 to id4 for other definitions.
4. *alnlen*: length of the query-target alignment (number of columns). The field is set to 0 if there is no alignment.
5. *mism*: number of mismatches in the alignment (zero or positive integer value).
6. *opens*: number of columns containing a gap opening (zero or positive integer value, excluding terminal gaps).
7. *qlo*: first nucleotide of the query aligned with the target. Always equal to 1 if there is an alignment, 0 otherwise (see *qilo* to ignore initial gaps).
8. *qhi*: last nucleotide of the query aligned with the target. Always equal to the length of the pairwise alignment, 0 otherwise (see *qihi* to ignore terminal gaps).
9. *tlo*: first nucleotide of the target aligned with the query. Always equal to 1 if there is an alignment, 0 otherwise (see *tilo* to ignore initial gaps).
10. *thi*: last nucleotide of the target aligned with the query. Always equal to the length of the pairwise alignment, 0 otherwise (see *tihi* to ignore terminal gaps).
11. *evalue*: expectancy-value (not computed for nucleotide alignments). Always set to -1.
12. *bits*: bit score (not computed for nucleotide alignments). Always set to 0.

--db *filename*

Compare query sequences (specified with --usearch_global) to the target sequences contained in *filename* in FASTA or FASTQ format, using global pairwise alignment. Alternatively, the name of a preformatted UDB database created using the makeudb_usearch command (see below) may be specified.

--dbmask *none|dust|soft*

Mask regions in the target database sequences using the dust method or the soft method, or do not mask (none). Warning, when using soft masking search commands become case sensitive. The default is to mask using dust.

--dbmatched *filename*

Write database target sequences matching at least one query sequence to *filename*, in fasta format. If the option --sizeout is used, the number of queries that matched each target sequence is indicated using the pattern ";size=*integer*";.

--dbnotmatched *filename*

Write database target sequences not matching query sequences to *filename*, in fasta format.

--fastapairs *filename*

Write pairwise alignments of query and target sequences to *filename*, in fasta format.

--fulldp Dummy option for compatibility with usearch. To maximize search sensitivity, **vsearch** uses a 8-way 16-bit SIMD vectorized full dynamic programming algorithm (Needleman-Wunsch), whether or not --fulldp is specified.

--gapext *string*

Set penalties for a gap extension. See --gapopen for a complete description of the penalty declaration system. The default is to initialize the six gap extending penalties using a penalty of 2 for extending internal gaps and a penalty of 1 for extending terminal gaps, in both query and target sequences (i.e. 2I/1E).

--gapopen *string*

Set penalties for a gap opening. A gap opening can occur in six different contexts: in the query (Q) or in the target (T) sequence, at the left (L) or right (R) extremity of the sequence, or inside the sequence (I). Sequence symbols (Q and T) can be combined with location symbols (L, I, and R), and numerical values to declare penalties for all possible contexts: aQL/bQI/cQR/dTL/eTI/fTR, where abcdef are zero or positive integers, and '/' is used as a separator.

To simplify declarations, the location symbols (L, I, and R) can be combined, the symbol (E) can be used to treat both extremities (L and R) equally, and the symbols Q and T can be omitted to treat query and target sequences equally. For instance, the default is to declare a penalty of 20 for opening internal gaps and a penalty of 2 for opening terminal gaps (left or right), in both query and target sequences (i.e. 20I/2E). If only a numerical value is given, without any sequence or location symbol, then the penalty applies to all gap openings. To forbid gap-opening, an infinite penalty value can be declared with the symbol '*'. To use **vsearch** as a semi-global aligner, a null-penalty can be applied to the left (L) or right (R) gaps.

vsearch always initializes the six gap opening penalties using the default parameters (20I/2E). The user is then free to declare only the values he/she wants to modify. The *string* is scanned from left to right, accepted symbols are (0123456789/LIREQT*), and later values override previous values.

Please note that **vsearch**, in contrast to usearch, only allows integer gap penalties. Because the lowest gap penalties are 0.5 by default in usearch, all default scores and gap penalties in **vsearch** have been doubled to maintain equivalent penalties and to produce identical alignments.

--hardmask

Mask sequence regions by replacing them with Ns instead of setting them to lower case as is the default. For more information, please see the Masking section.

--id *real* Reject the sequence match if the pairwise identity is lower than *real* (value ranging from 0.0 to 1.0 included). The search process sorts target sequences by decreasing number of *k*-mers they have in common with the query sequence, using that information as a proxy for sequence similarity. That efficient pre-filtering also prevents pairwise alignments with very short, or with weakly matching targets, as there needs to be by default at least 12 shared *k*-mers to start the pairwise alignment, and at least one out of every 16 *k*-mers from the query needs to match the target (see options --wordlength and --minwordmatches to change that behaviour). Consequently, using values lower than --id 0.5 is not likely to capture more weakly matching targets. The pairwise identity is by default defined as the number of (matching columns) / (alignment length - terminal gaps). That definition can be modified by --iddef.

--iddef 0|1|2|3|4

Change the pairwise identity definition used in --id. Values accepted are:

0. CD-HIT definition: (matching columns) / (shortest sequence length).
1. edit distance: (matching columns) / (alignment length).
2. edit distance excluding terminal gaps (default definition for --id).
3. Marine Biological Lab definition counting each gap opening (internal or terminal) as a single mismatch, whether or not the gap was extended: $1.0 - [(mismatches + gap openings)/(longest sequence length)]$
4. BLAST definition, equivalent to --iddef 1 for global pairwise alignments.

The option --userfields accepts the fields id0 to id4, in addition to the field id, to report the pairwise identity values corresponding to the different definitions.

--idprefix *positive integer*

Reject the sequence match if the first *integer* nucleotides of the target do not match the query.

--idsuffix *positive integer*

Reject the sequence match if the last *integer* nucleotides of the target do not match the query.

--lca_cutoff *real*

Adjust the fraction of matching hits required for the last common ancestor (LCA) output with the --lcaout option during searches. The default value is 1.0 which requires all hits to match at each taxonomic rank for that rank to be included. If a lower cutoff value is used, e.g. 0.95, a small fraction of non-matching hits are allowed while that rank will still be reported. The argument to this option must be larger than 0.5, but not larger than 1.0.

--lcaout *filename*

Output last common ancestor (LCA) information about the hits of each query to a text file in a tab-separated format. The first column contains the query id, while the second column contains the taxonomic information. The headers of the sequences in the database must contain taxonomic information in the same format as used with the --syntax command, e.g. "tax=k:Archaea,p:Euryarchaeota,c:Halobacteria". Only the initial parts of the taxonomy that are common to a large fraction of the hits of each query will be output. It is necessary to set the --maxaccepts option to a value different from 1 for this information to be useful. The --top_hits_only option may also be useful. The fraction of matching hits required may be adjusted by the --lca_cutoff option (default 1.0).

--leftjust Reject the sequence match if the pairwise alignment begins with gaps.**--lengthout**

Write sequence length information to the output files in FASTA format by adding a ";length=*integer*" attribute in the header.

--match *integer*

Score assigned to a match (i.e. identical nucleotides) in the pairwise alignment. The default value is 2.

--matched *filename*

Write query sequences matching database target sequences to *filename*, in fasta format.

--maxaccepts *positive integer*

Maximum number of matching target sequences to accept before stopping the search for a given query. The default value is 1. This option works in pair with --maxrejects. The search process sorts target sequences by decreasing number of *k*-mers they have in common with the query sequence, using that information as a proxy for sequence

similarity. After pairwise alignments, if the first target sequence passes the acceptance criteria, it is accepted as best hit and the search process stops for that query. If `--maxaccepts` is set to a higher value, more matching targets are accepted. If `--maxaccepts` and `--maxrejects` are both set to 0, the complete database is searched. See `--maxhits` option for a control on the number of hits reported per query when search is done on both strands.

--maxdiffs *positive integer*

Reject the sequence match if the alignment contains more than *integer* substitutions, insertions or deletions.

--maxgaps *positive integer*

Reject the sequence match if the alignment contains more than *integer* insertions or deletions.

--maxhits *non-negative integer*

Maximum number of hits to show once the search is terminated for a given query (hits are sorted by decreasing identity). When searching only on the plus strand (default situation, see `--strand`), the number of matching targets (`--maxaccepts`) and the number of hits (`--maxhits`) are the same. However, when searching on both strands, there could be two hits per target (one per strand): `--maxhits` then controls the overall number of reported hits per query. Unlimited by default or if the argument is zero. This option applies to `--alnout`, `--blast6out`, `--fastapairs`, `--samout`, `--uc`, or `--userout` output files.

--maxid *real*

Reject the sequence match if the percentage of identity between the two sequences is greater than *real*.

--maxqsize *positive integer*

Reject query sequences with an abundance greater than *integer*.

--maxqt *real*

Reject if the query/target sequence length ratio is greater than *real*.

--maxrejects *positive integer*

Maximum number of non-matching target sequences to consider before stopping the search for a given query. The default value is 32. This option works in pair with `--maxaccepts`. The search process sorts target sequences by decreasing number of *k*-mers they have in common with the query sequence, using that information as a proxy for sequence similarity. After pairwise alignments, if none of the first 32 examined target sequences pass the acceptance criteria, the search process stops for that query (no hit). If `--maxrejects` is set to a higher value, more target sequences are considered. If `--maxaccepts` and `--maxrejects` are both set to 0, the complete database is searched.

--maxsizeratio *real*

Reject if the query/target abundance ratio is greater than *real*.

--maxsl *real*

Reject if the shorter/longer sequence length ratio is greater than *real*.

--maxsubs *positive integer*

Reject the sequence match if the pairwise alignment contains more than *integer* substitutions.

--mid *real*

Reject the sequence match if the percentage of identity is lower than *real* (ignoring all gaps, internal and terminal).

--mincols *positive integer*

Reject the sequence match if the alignment length is shorter than *integer*.

- minqt *real***
Reject if the query/target sequence length ratio is lower than *real*.
- minsizeratio *real***
Reject if the query/target abundance ratio is lower than *real*.
- minsl *real***
Reject if the shorter/longer sequence length ratio is lower than *real*.
- mintsize *positive integer***
Reject target sequences with an abundance lower than *integer*.
- minwordmatches *non-negative integer***
Minimum number of *k*-mers or word matches required for a sequence to be considered further. Default value is 12 for the default word length 8. For word lengths 3-15, the default minimum word matches are 18, 17, 16, 15, 14, 12, 11, 10, 9, 8, 7, 5 and 3, respectively. If the query sequence has fewer unique words than the number specified, all words in the query must match. If the argument is 0, no word matches are required.
- mismatch *integer***
Score assigned to a mismatch (i.e. different nucleotides) in the pairwise alignment. The default value is -4.
- mothur_shared_out *filename***
Write search results to an OTU table in the mothur 'shared' tab-separated plain text file format. The query file contains the samples, while the database file contains the OTUs. Sample and OTU identifiers are extracted from the header of these sequences. See the --otutabout option in the Clustering section for further details.
- notmatched *filename***
Write query sequences not matching database target sequences to *filename*, in fasta format.
- otutabout *filename***
Write search results to an OTU table in the classic tab-separated plain text format. The query file contains the samples, while the database file contains the OTUs. Sample and OTU identifiers are extracted from the header of these sequences (--sample option). See the --mothur_shared_out option in the Clustering section for further details.
- output_no_hits**
Write both matching and non-matching queries to --alnout, --blast6out, --samout or --userout output files. Non-matching queries are labelled 'No hits' in --alnout files.
- pattern *string***
This option is ignored. It is provided for compatibility with usearch.
- qmask *none|dust|soft***
Mask regions in the query sequences using the dust or the soft algorithms, or do not mask (none). Warning, when using soft masking search commands become case sensitive. The default is to mask using *dust*.
- qsegout *filename***
Write the aligned part of each query sequence to *filename* in FASTA format.
- query_cov *real***
Reject if the fraction of the query aligned to the target sequence is lower than *real* (value ranging from 0.0 to 1.0 included). The query coverage is computed as (matches + mismatches) / query sequence length. Internal or terminal gaps are not taken into account.
- rightjust**
Reject the sequence match if the pairwise alignment ends with gaps.

--rowlen *positive integer*

Width of alignment lines in --alnout output. The default value is 64. Set to 0 to eliminate wrapping.

--samheader

Include header lines to the SAM file when --samout is specified. The header includes lines starting with @HD, @SQ and @PG, but no @RG lines (see ([link](#)) <<https://github.com/samtools/hts-specs>> <<https://github.com/samtools/hts-specs>>). By default no header line is written.

--samout *filename*

Write alignment results to *filename* using the SAM format (a tab-separated text file). When using the --samheader option, the SAM file starts with header lines. Each non-header line is a SAM record, which represents either a query-target alignment or the absence of match for a query (output order may vary when using multiple threads). Each record contains 11 mandatory fields and optional fields (see ([link](#)) <<https://github.com/samtools/hts-specs>> <<https://github.com/samtools/hts-specs>> for a complete description of the format):

1. query sequence label.
2. combination of bitwise flags. Possible values are: 0 (top hit), 4 (no hit), 16 (reverse-complemented hit), 256 (secondary hit, i.e. all hits except the top hit).
3. target sequence label.
4. first position of a target aligned with the query (always 1 for global pairwise alignments, 0 if there is no match).
5. mapping quality (ignored, always set to '*').
6. CIGAR string (set to '*' if there is no match).
7. name of the target sequence matching with the next read of the query (for mate reads only, ignored and always set to '*').
8. position of the primary alignment of the next read of the query (for mate reads only, ignored and always set to 0).
9. target sequence length (for multi-segment targets, ignored and always set to 0).
10. query sequence (complete, not only the segment aligned to the target as usearch does).
11. quality string (ignored, always set to '*').

Optional fields for query-target matches (number and order of fields may vary):

12. AS:i:? alignment score (i.e. percentage of identity).
13. XN:i:? next best alignment score (always set to 0).
14. XM:i:? number of mismatches.
15. XO:i:? number of gap openings (excluding terminal gaps).
16. XG:i:? number of gap extensions (excluding terminal gaps).
17. NM:i:? edit distance to the target (sum of XM and XG).
18. MD:Z:? string for mismatching positions.
19. YT:Z:UU string representing the alignment type.

--search_exact *filename*

Search for exact full-length matches to the query sequences contained in *filename* in the fasta database of target sequences (--db). Only 100% exact matches are reported and this command is much faster than --usearch_global. The --id, --maxaccepts and --maxrejects options are ignored, but the rest of the searching options may be specified.

--self Reject the sequence match if the query and target labels are identical.

--selfid Reject the sequence match if the query and target sequences are strictly identical.

--sizeout Add abundance annotations to the output of the option --dbmatched (using the pattern ';size=integer;'), to report the number of queries that matched each target.

--strand *plus|both*

When searching for similar sequences, check the *plus* strand only (default) or check *both* strands.

--target_cov *real*

Reject the sequence match if the fraction of the target sequence aligned to the query sequence is lower than *real*. The target coverage is computed as (matches + mismatches) / target sequence length. Internal or terminal gaps are not taken into account.

--top_hits_only

Only the top hits with an equally high percentage of identity between the query and database sequence sets are written to the output specified with the options --lcaout, --alnout, --samout, --userout, --blast6out, --uc, --fastapairs, --matched or --notmatched (but not --dbmatched and --dbnotmatched). For each query, the top hit is the one presenting the highest percentage of identity (see the --iddef option to change the way identity is measured). For a given query, if several top hits present exactly the same percentage of identity, the number of matching targets reported is controlled by the --maxaccepts value (1 by default), and the number of hits is controlled by the --maxhits option.

--tsegout *filename*

Write the aligned part of each target sequence to *filename* in FASTA format.

--uc *filename*

Output searching results in *filename* using a tab-separated uclust-like format with 10 columns. When using the --search_exact command, the table layout is the same than with the --allpairs_global. When using the --usearch_global command, the table present two different type of entries: hit (H) or no hit (N). Each query sequence is compared to all other sequences, and the best hit (--maxaccepts 1) or several hits (--maxaccepts > 1 and --uc_allhits) are reported (H). Output order may vary when using multiple threads. Column content varies with the type of entry (H or N):

1. Record type: H, or N ('hit' or 'no hit').
2. Ordinal number of the target sequence (based on input order, starting from zero). Set to '*' for N.
3. Sequence length. Set to '*' for N.
4. Percentage of similarity with the target sequence. Set to '*' for N.
5. Match orientation + or -. Set to '?' for N.
6. Not used, always set to zero for H, or '*' for N.
7. Not used, always set to zero for H, or '*' for N.
8. Compact representation of the pairwise alignment using the CIGAR format (Compact Idiosyncratic Gapped Alignment Report): M (match/mismatch), D (deletion) and I (insertion). The equal sign '=' indicates that the query is identical to the centroid sequence (ignoring terminal gaps). Set to '*' for N.

9. Label of the query sequence.
10. Label of the target centroid sequence. Set to '*' for N.

--uc_allhits

With the commands --search_exact and --usearch_global, when using the --uc option, show all hits, not just the top hit for each query.

--usearch_global filename

Compare target sequences (--db) to the query sequences contained in *filename* in FASTA or FASTQ format, using global pairwise alignment.

--userfields string

When using --userout, select and order the fields written to the output file. Fields are separated by '+' (e.g. query+target+id). See the 'Userfields' section for a complete list of fields.

--userout filename

Write user-defined tab-separated output to *filename*. Select the fields with the option --userfields. Output order may vary when using multiple threads. If --userfields is empty or not present, *filename* is empty.

--weak_id real

Show hits with percentage of identity of at least *real*, without terminating the search. A normal search stops as soon as enough hits are found (as defined by --maxaccepts, --maxrejects, and --id). As --weak_id reports weak hits that are not deduced from --maxaccepts (but count towards --maxrejects), high --id values can be used, hence preserving both speed and sensitivity. Logically, *real* must be smaller than the value indicated by --id.

--wordlength positive integer

Length of words (i.e. *k*-mers) for database indexing. The range of possible values goes from 3 to 15, but values near 8 or 9 are generally recommended. Longer words may reduce the sensitivity/recall for weak similarities, but can increase precision. On the other hand, shorter words may increase sensitivity or recall, but may reduce precision. Computation time generally increases with shorter words and decreases with longer words, but it increases again for very long words. Memory requirements for a part of the index increase with a factor of 4 each time word length increases by one nucleotide, and this generally becomes significant for long words (12 or more). The default value is 8.

--xlength Strip header attribute ";length=*integer*" from input sequences. This attribute is added to output sequences by the --lengthout option.

Shuffling options:

Fasta entries in the input file are outputted in a pseudo-random order.

--lengthout

Write sequence length information to the output files in FASTA format by adding a ";length=*integer*" attribute in the header.

--output filename

Write the shuffled sequences to *filename*, in fasta format.

--randseed integer

When shuffling sequence order, use *integer* as seed. A given seed always produces the same output order (useful for replicability). Set to 0 to use a pseudo-random seed (default behaviour).

--relabel string

Relabel sequences using the prefix *string* and a ticker (1, 2, 3, etc.) to construct the new headers. Use --sizeout to conserve the abundance annotations.

--relabel_keep

When relabelling, keep the old identifier in the header after a space.

--relabel_md5

Relabel sequences using the MD5 message digest algorithm applied to each sequence. Former sequence headers are discarded. The sequence is converted to upper case and U is replaced by T before the digest is computed. The MD5 digest is a cryptographic hash function designed to minimize the probability that two different inputs gives the same output, even for very similar, but non-identical inputs. Still, there is always a very small, but non-zero probability that two different inputs give the same result. The MD5 digest generates a 128-bit (16-byte) digest that is represented by 16 hexadecimal numbers (using 32 symbols among 0123456789abcdef). Use --sizeout to conserve the abundance annotations.

--relabel_self

Relabel sequences using the sequence itself as the label.

--relabel_sha1

Relabel sequences using the SHA1 message digest algorithm applied to each sequence. It is similar to the --relabel_md5 option but uses the SHA1 algorithm instead of the MD5 algorithm. The SHA1 digest generates a 160-bit (20-byte) result that is represented by 20 hexadecimal numbers (40 symbols). The probability of a collision (two non-identical sequences having the same digest) is smaller for the SHA1 algorithm than it is for the MD5 algorithm. Use --sizeout to conserve the abundance annotations.

--sizeout When using --relabel, --relabel_self, --relabel_md5 or --relabel_sha1, preserve and report abundance annotations to the output fasta file (using the pattern '`;size=integer;`').

--shuffle *filename*

Pseudo-randomly shuffle the order of sequences contained in *filename*.

--topn *positive integer*

Output only the first *integer* sequences after pseudo-random reordering.

--xlength Strip header attribute "`;length=integer`" from input sequences. This attribute is added to output sequences by the --lengthout option.

--xsize Strip abundance information from the headers when writing the output file.

Sorting options:

Fasta entries are sorted by decreasing abundance (--sortbysize) or sequence length (--sortbylength). To obtain a stable sorting order, ties are sorted by decreasing abundance (if present) and label increasing alpha-numerical order (--sortbylength), or just by label increasing alpha-numerical order (--sortbysize). Label sorting assumes that all sequences have unique labels. The same applies to the automatic sorting performed during chimera checking (--uchime_denovo), dereplication (--derep_fulllength), and clustering (--cluster_fast and --cluster_size).

--lengthout

Write sequence length information to the output files in FASTA format by adding a "`;length=integer`" attribute in the header.

--maxsize *positive integer*

When using --sortbysize, discard sequences with an abundance value greater than *integer*.

--minsize *positive integer*

When using --sortbysize, discard sequences with an abundance value smaller than *integer*.

--output *filename*

Write the sorted sequences to *filename*, in fasta format.

--relabel *string*

Please see the description of the same option under Chimera detection for details.

--relabel_keep

When relabelling, keep the old identifier in the header after a space.

--relabel_md5

Please see the description of the same option under Chimera detection for details.

--relabel_self

Please see the description of the same option under Chimera detection for details.

--relabel_sha1

Please see the description of the same option under Chimera detection for details.

--sizeout When using --relabel, report abundance annotations to the output fasta file (using the pattern ';size=*integer*;').

--sortbylength *filename*

Sort by decreasing length the sequences contained in *filename*. See the general options --minseqlength and --maxseqlength to eliminate short and long sequences.

--sortbysize *filename*

Sort by decreasing abundance the sequences contained in *filename* (missing abundance values are assumed to be ';size=1'). See the options --minsize and --maxsize to eliminate rare and dominant sequences.

--topn *positive integer*

Output only the top *integer* sequences (i.e. the longest or the most abundant).

--xlength Strip header attribute ";length=*integer*" from input sequences. This attribute is added to output sequences by the --lengthout option.

--xsize Strip abundance information from the headers when writing the output file.

Subsampling options:

Subsampling randomly extracts a certain number or a certain percentage of the sequences in the input file. If the --sizein option is in effect, the abundances of the input sequences is taken into account and the sampling is performed as if the input sequences were rereplicated, subsampled and dereplicated before being written to the output file. The extraction is performed as a random sampling with a uniform distribution among the input sequences and is performed without replacement. The input file is specified with the --fastx_subsample option, the output files are specified with the --fastaout and --fastqout options and the amount of sequences to be sampled is specified with the --sample_pct or --sample_size options. The sequences not sampled may be written to files specified with the options --fasta_discarded and --fastq_discarded. The --fastq_ascii, --fastq_qmin and --fastq_qmax options are also available.

--fastaout *filename*

Write the sampled sequences to *filename*, in fasta format.

--fastaout_discarded *filename*

Write the sequences not sampled to *filename*, in fasta format.

--fastq_ascii *positive integer*

Define the ASCII character number used as the basis for the FASTQ quality score. The default is 33, which is used by the Sanger / Illumina 1.8+ FASTQ format (phred+33).

The value 64 is used by the Solexa, Illumina 1.3+ and Illumina 1.5+ formats (phred+64). Only 33 and 64 are valid arguments.

--fastq_qmax *positive integer*

Specify the maximum quality score accepted when reading FASTQ files. The default is 41, which is usual for recent Sanger/Illumina 1.8+ files.

--fastq_qmin *positive integer*

Specify the minimum quality score accepted for FASTQ files. The default is 0, which is usual for recent Sanger/Illumina 1.8+ files. Older formats may use scores between -5 and 2.

--fastqout *filename*

Write the sampled sequences to *filename*, in fastq format. Requires input in fastq format.

--fastqout_discarded *filename*

Write the sequences not sampled to *filename*, in fastq format. Requires input in fastq format.

--fastx_subsample *filename*

Perform subsampling from the sequences in the specified input file that is in FASTA or FASTQ format.

--lengthout

Write sequence length information to the output files in FASTA format by adding a ";length=*integer*" attribute in the header.

--randseed *positive integer*

When subsampling, use *integer* as a seed for the pseudo-random generator. A given seed always produces the same output, which is useful for replicability. Set to 0 to use a pseudo-random seed (default behaviour).

--relabel *string*

Relabel sequences using the prefix *string* and a ticker (1, 2, 3, etc.) to construct the new headers. Use --sizeout to conserve the abundance annotations.

--relabel_keep

When relabelling, keep the old identifier in the header after a space.

--relabel_md5

Relabel sequences using the MD5 message digest algorithm applied to each sequence. Former sequence headers are discarded. The sequence is converted to upper case and U is replaced by T before the digest is computed. The MD5 digest is a cryptographic hash function designed to minimize the probability that two different inputs give the same output, even for very similar, but non-identical inputs. Still, there is always a very small, but non-zero probability that two different inputs give the same result. The MD5 digest generates a 128-bit (16-byte) digest that is represented by 16 hexadecimal numbers (using 32 symbols among 0123456789abcdef). Use --sizeout to conserve the abundance annotations.

--relabel_self

Relabel sequences using the sequence itself as the label.

--relabel_sha1

Relabel sequences using the SHA1 message digest algorithm applied to each sequence. It is similar to the --relabel_md5 option but uses the SHA1 algorithm instead of the MD5 algorithm. The SHA1 digest generates a 160-bit (20-byte) result that is represented by 20 hexadecimal numbers (40 symbols). The probability of a collision (two non-identical sequences having the same digest) is smaller for the SHA1 algorithm than it is for the MD5 algorithm. Use --sizeout to conserve the abundance annotations.

--sample_pct *real*

Subsample the given percentage of the input sequences. Accepted values range from 0.0 to 100.0.

--sample_size *positive integer*

Extract the given number of sequences.

- sizein** Take the abundance information of the input file into account, otherwise the abundance of each sequence is considered to be 1.
- sizeout** Write abundance information to the output file.
- xlength** Strip header attribute ";length=integer" from input sequences. This attribute is added to output sequences by the --lengthout option.
- xsize** Strip abundance information from the headers when writing the output file.

Taxonomic classification options:

The vsearch command --sintax will classify the input sequences according to the Sintax algorithm as described by Robert Edgar (2016) in SINTAX: a simple non-Bayesian taxonomy classifier for 16S and ITS sequences, BioRxiv, 074161. Preprint. doi: 10.1101/074161 ([link](#)) (<https://doi.org/10.1101/074161>)

The name of the fasta file containing the input sequences to be classified is given as an argument to the --sintax command. The reference sequence database is specified with the --db option. The results are written in a tab delimited text file whose name is specified with the --tabbedout option. The --sintax_cutoff option may be used to set a minimum level of bootstrap support for the taxonomic ranks to be reported. The --randseed option may be included to specify a seed for initialisation of the random number generator used by the algorithm. Please note that when using multiple threads, the --randseed option may not work as intended, because sequences may be processed in a random order by different threads. To ensure the same results each time, use a single thread --threads 1) in combination with a fixed random seed specified with --randseed.

Multithreading is supported. Databases in UDB files are supported. The strand option may be specified.

The reference database must contain taxonomic information in the header of each sequence in the form of a string starting with ";tax=" and followed by a comma-separated list of up to nine taxonomic identifiers. Each taxonomic identifier must start with an indication of the rank by one of the letters d (for domain) k (kingdom), p (phylum), c (class), o (order), f (family), g (genus), s (species), or t (strain). The letter is followed by a colon (:) and the name of that rank. Commas and semicolons are not allowed in the name of the rank. Non-ascii characters should be avoided in the names.

Example:

```
>X80725_S000004313;tax=d:Bacteria,p:Proteobacteria,c:Gammaproteobacteria,o:Enterobacteriales,f:Enterobacteriaceae,g:Escherichia/Shigella,s:Escherichia_coli,t:str._K-12_substr._MG1655
```

The option --notrunclabels is turned on by default for this command, allowing spaces in the taxonomic identifiers.

If two sequences in the reference database has equally many kmer matches with the query, the shortest sequence will be chosen by default. If they are equally long, the sequence appearing first in the database will be chosen. If the recommended option --sintax_random is specified, sequences with an equal number of kmer matches will instead be chosen by a random draw.

--db *filename*

Read the reference sequences from *filename*, in FASTA, FASTQ or UDB format. These sequences need to be annotated with taxonomy.

--randseed *positive integer*

Use *integer* as seed for the random number generator used in the Sintax algorithm. A given seed always produces the same output order (useful for replicability). Set to 0 to use a pseudo-random seed (default behaviour). Does not work correctly with multiple threads; please use --threads 1 to ensure correct behaviour.

--syntax *filename*

Read the input sequences from *filename*, in FASTA or FASTQ format.

--syntax_cutoff *real*

Specify a minimum level of bootstrap support for the taxonomic ranks that will be included in column 4 of the output file. For instance 0.9, corresponding to 90%.

--syntax_random

Break ties between sequences with equally many kmer matches by a random draw. This option is recommended and may be made the default in the future.

--tabbedout *filename*

Write the results to *filename*, in a tab-separated text format. Column 1 contains the query label. Column 2 contains the predicted taxonomy in the same format as for the reference data, with bootstrap support indicated in parentheses after each rank. Column 3 contains the strand. If the --syntax_cutoff option is used, the predicted taxonomy will be repeated in column 4 while omitting the bootstrap values and including only the ranks with support at or above the threshold.

UDB options:

Databases to be used with the --usearch_global or --syntax commands may be prepared from FASTA files and stored to a binary UDB formatted file in order to speed up searching. This may be worthwhile when searching a large database repeatedly. The sequences are indexed and stored in a way that can be quickly loaded into memory. The commands and options below can be used to create and inspect UDB files.

--dbmask *none|dust|soft*

Specify the sequence masking method used with the --makeudb_usearch command, either none, dust or soft. No masking is performed when none is specified. When dust is specified, the DUST algorithm will be used for masking low complexity regions (short repeats and skewed composition). Lower case letters in the input file will be masked when soft is specified (soft masking).

--hardmask

Mask sequences by replacing letters with N for the --makeudb_usearch command. The default is to use lower case letters (soft masking).

--makeudb_usearch *filename*

Create an UDB database file from the FASTA-formatted sequences in the file with the given *filename*. The UDB database is written to the file specified with the --output option.

--output *filename*

Specify the *filename* of a FASTA or UDB output file for the --makeudb_usearch or the --fdb2fasta command, respectively.

--fdb2fasta *filename*

Read the UDB database in the file with the given *filename* and output the sequences in FASTA format in the file specified by the --output option.

--fdbinfo *filename*

Show information about the UDB database in the file with the given *filename*.

--fdbstats *filename*

Report statistics about the indexed words in the UDB database in the file with the given *filename*.

--wordlength *positive integer*

Specify the length of the words to be used when creating the UDB database index using the --makeudb_usearch command. Valid numbers range from 3 to 15. The default is 8.

Userfields (fields accepted by the --userfields option):

aln	Print a string of M (match/mismatch, i.e. not a gap), D (delete, i.e. a gap in the query) and I (insert, i.e. a gap in the target) representing the pairwise alignment. Empty field if there is no alignment.
alnlen	Print the length of the query-target alignment (number of columns). The field is set to 0 if there is no alignment.
bits	Bit score (not computed for nucleotide alignments). Always set to 0.
caln	Compact representation of the pairwise alignment using the CIGAR format (Compact Idiosyncratic Gapped Alignment Report): M (match/mismatch), D (deletion) and I (insertion). The equal sign '=' indicates that the query is identical to the centroid sequence (ignoring terminal gaps). Empty field if there is no alignment.
evalue	E-value (not computed for nucleotide alignments). Always set to -1.
exts	Number of columns containing a gap extension (zero or positive integer value).
gaps	Number of columns containing a gap (zero or positive integer value, excluding terminal gaps).
id	The percentage of identity, according to the identity definition specified by the --iddef option. Equal to id0, id1, id2, id3 or id4 below. By default the same as id2.
id0	CD-HIT definition of the percentage of identity (real value ranging from 0.0 to 100.0) using the length of the shortest sequence in the pairwise alignment as denominator: $100 * (\text{matching columns}) / (\text{shortest sequence length})$.
id1	The percentage of identity (real value ranging from 0.0 to 100.0) is defined as the edit distance: $100 * (\text{matching columns}) / (\text{alignment length})$.
id2	The percentage of identity (real value ranging from 0.0 to 100.0) is defined as the edit distance, excluding terminal gaps.
id3	Marine Biological Lab definition of the percentage of identity (real value ranging from 0.0 to 100.0), counting each gap opening (internal or terminal) as a single mismatch, whether or not the gap was extended, and using the length of the longest sequence in the pairwise alignment as denominator: $100 * (1.0 - [(\text{mismatches} + \text{gaps}) / (\text{longest sequence length})])$.
id4	BLAST definition of the percentage of identity (real value ranging from 0.0 to 100.0), equivalent to --iddef 1 in a context of global pairwise alignment. The field id4 is always equal to the field id1.
ids	Number of matches in the alignment (zero or positive integer value).
mism	Number of mismatches in the alignment (zero or positive integer value).
opens	Number of columns containing a gap opening (zero or positive integer value, excluding terminal gaps).
pairs	Number of columns containing only nucleotides. That value corresponds to the length of the alignment minus the gap-containing columns (zero or positive integer value).
pctgaps	Number of columns containing gaps expressed as a percentage of the alignment length (real value ranging from 0.0 to 100.0).
pctpv	Percentage of positive columns. When working with nucleotide sequences, this is equivalent to the percentage of matches (real value ranging from 0.0 to 100.0).
pv	Number of positive columns. When working with nucleotide sequences, this is equivalent to the number of matches (zero or positive integer value).

qcov	Fraction of the query sequence that is aligned with the target sequence (real value ranging from 0.0 to 100.0). The query coverage is computed as $100.0 * (\text{matches} + \text{mismatches}) / \text{query sequence length}$. Internal or terminal gaps are not taken into account. The field is set to 0.0 if there is no alignment.
qframe	Query frame (-3 to +3). That field only concerns coding sequences and is not computed by vsearch . Always set to +0.
qli	Last nucleotide of the query aligned with the target. Always equal to the length of the pairwise alignment, 0 otherwise (see <i>qli</i> to ignore terminal gaps).
qlio	Last nucleotide of the query aligned with the target (ignoring terminal gaps). Nucleotide numbering starts from 1. The field is set to 0 if there is no alignment.
ql	Query sequence length (positive integer value). The field is set to 0 if there is no alignment.
qlo	First nucleotide of the query aligned with the target. Always equal to 1 if there is an alignment, 0 otherwise (see <i>qlo</i> to ignore initial gaps).
qrow	Print the sequence of the query segment as seen in the pairwise alignment (i.e. with gap insertions if need be). Empty field if there is no alignment.
qs	Query segment length. Always equal to query sequence length.
qstrand	Query strand orientation (+ or - for nucleotide sequences). Empty field if there is no alignment.
query	Query label.
raw	Raw alignment score (negative, null or positive integer value). The score is the sum of match rewards minus mismatch penalties, gap openings and gap extensions. The field is set to 0 if there is no alignment.
target	Target label. The field is set to '*' if there is no alignment.
tcov	Fraction of the target sequence that is aligned with the query sequence (real value ranging from 0.0 to 100.0). The target coverage is computed as $100.0 * (\text{matches} + \text{mismatches}) / \text{target sequence length}$. Internal or terminal gaps are not taken into account. The field is set to 0.0 if there is no alignment.
tframe	Target frame (-3 to +3). That field only concerns coding sequences and is not computed by vsearch . Always set to +0.
thi	Last nucleotide of the target aligned with the query. Always equal to the length of the pairwise alignment, 0 otherwise (see <i>thi</i> to ignore terminal gaps).
tih	Last nucleotide of the target aligned with the query (ignoring terminal gaps). Nucleotide numbering starts from 1. The field is set to 0 if there is no alignment.
tilo	First nucleotide of the target aligned with the query (ignoring initial gaps). Nucleotide numbering starts from 1. The field is set to 0 if there is no alignment.
tl	Target sequence length (positive integer value). The field is set to 0 if there is no alignment.
tlo	First nucleotide of the target aligned with the query. Always equal to 1 if there is an alignment, 0 otherwise (see <i>tlo</i> to ignore initial gaps).
trow	Print the sequence of the target segment as seen in the pairwise alignment (i.e. with gap insertions if need be). Empty field if there is no alignment.

- ts** Target segment length. Always equal to target sequence length. The field is set to 0 if there is no alignment.
- tstrand** Target strand orientation (+ or - for nucleotide sequences). Always set to '+', so reverse strand matches have tstrand '+' and qstrand '-'. Empty field if there is no alignment.

DELIBERATE CHANGES

If you are a usearch user, our objective is to make you feel at home. That's why **vsearch** was designed to behave like usearch, to some extent. Like any complex software, usearch is not free from quirks and inconsistencies. We decided not to reproduce some of them, and for complete transparency, to document here the deliberate changes we made.

During a search with usearch, when using the options --blast6out and --output_no_hits, for queries with no match the number of fields reported is 13, where it should be 12. This is corrected in **vsearch**.

The field raw of the --userfields option is not informative in usearch. This is corrected in **vsearch**.

The fields qlo, qhi, tlo, thi now have counterparts (qilo, qih, tilo, tihi) reporting alignment coordinates ignoring terminal gaps.

In usearch, when using the option --output_no_hits, queries that receive no match are reported in --blast6out file, but not in the alignment output file. This is corrected in **vsearch**.

vsearch introduces a new --cluster_size command that sorts sequences by decreasing abundance before clustering.

vsearch reintroduces --iddef alternative pairwise identity definitions that were removed from usearch.

vsearch extends the --topn option to sorting commands.

vsearch extends the --sizein option to dereplication (--deref_fulllength) and clustering (--cluster_fast).

vsearch treats T and U as identical nucleotides during dereplication.

vsearch sorting is stabilized by using sequence abundances or sequences labels as secondary or tertiary keys.

vsearch by default uses the DUST algorithm for masking low-complexity regions. Masking behaviour is also slightly changed to be more consistent.

NOVELTIES

vsearch introduces new commands and new options not present in usearch 7. They are described in the 'Options' section of this manual. Here is a short list:

- uchime2_denovo, uchime3_denovo, alignwidth, borderline, fasta_score (chimera checking)
- cluster_size, cluster_unequal, clusterout_id, clusterout_sort, profile (clustering)
- fasta_width, gzip_decompress, bzip2_decompress (general option)
- iddef (clustering, pairwise alignment, searching)
- maxuniquesize (dereplication)
- relabel_md5, relabel_self and relabel_sha1 (chimera detection, dereplication, FASTQ processing, shuffling, sorting)
- shuffle (shuffling)
- fastq_eestats, fastq_eestats2, fastq_maxlen, fastq_trunc (FASTQ processing)
- fastqout_discarded, fastqout_discarded (subsampling)
- rereplicate (dereplication/rereplication)

EXAMPLES

Align all sequences in a database with each other and output all pairwise alignments:

```
vsearch --allpairs_global database.fas --alnout results.aln --acceptall
```

Check for the presence of chimeras (*de novo*); parents should be at least 1.5 times more abundant than

chimeras. Output non-chimeric sequences in fasta format (no wrapping):

```
vsearch --uchime_denovo queries.fas --abskew 1.5 --nonchimeras results.fas --fasta_width 0
```

Cluster with a 97% similarity threshold, collect cluster centroids, and write cluster descriptions using a uclust-like format:

```
vsearch --cluster_fast queries.fas --id 0.97 --centroids centroids.fas --uc clusters.uc
```

Dereplicate the sequences contained in *queries.fas*, take into account the abundance information already present, write unwrapped fasta sequences to *queries_unique.fas* with the new abundance information, discard all sequences with an abundance of 1:

```
vsearch --deref_fulllength queries.fas --sizein --fasta_width 0 --sizeout --output queries_unique.fas --minuniquesize 2
```

Mask simple repeats and low complexity regions in the input fasta file with the DUST algorithm (masked regions are lowercased), and write the results to the output file:

```
vsearch --maskfasta queries.fas --qmask dust --output queries_masked.fas
```

Search queries in a reference database, with a 80%-similarity threshold, take terminal gaps into account when calculating pairwise similarities, output pairwise alignments:

```
vsearch --usearch_global queries.fas --db references.fas --id 0.8 --iddef 1 --alnout results.aln
```

Search a sequence dataset against itself (ignore self hits), get all matches with at least 60% similarity, and collect results in a blast-like tab-separated format. Accept an unlimited number of hits (--maxaccepts 0), and compare each query to all other sequences, including unlikely candidates (--maxrejects 0):

```
vsearch --usearch_global queries.fas --db queries.fas --self --id 0.6 --blast6out results.blast6 --maxaccepts 0 --maxrejects 0
```

Shuffle the input fasta file (change the order of sequences) in a repeatable fashion (fixed seed), and write unwrapped fasta sequences to the output file:

```
vsearch --shuffle queries.fas --output queries_shuffled.fas --randseed 13 --fasta_width 0
```

Sort by decreasing abundance the sequences contained in *queries.fas* (using the 'size=integer' information), relabel the sequences while preserving the abundance information (with --sizeout), keep only sequences with an abundance equal to or greater than 2:

```
vsearch --sortysize queries.fas --output queries_sorted.fas --relabel sampleA_ --sizeout --min-size 2
```

AUTHORS

Implementation and documentation by Torbjørn Rognes, Frédéric Mahé and Tomás Flouri.

CITATION

Rognes T, Flouri T, Nichols B, Quince C, Mahé F. (2016) VSEARCH: a versatile open source tool for metagenomics. *PeerJ* 4:e2584 doi: 10.7717/peerj.2584 ([link](#)) (<https://doi.org/10.7717/peerj.2584>)

REPORTING BUGS

Submit suggestions and bug-reports at ([link](#)) (<https://github.com/torognes/vsearch/issues>) <<https://github.com/torognes/vsearch/issues>>, send a pull request on ([link](#)) (<https://github.com/torognes/vsearch>) <<https://github.com/torognes/vsearch>>, or compose a friendly or curmudgeon e-mail to Torbjørn Rognes ([link](#)) (torognes@ifi.uio.no) <torognes@ifi.uio.no>.

AVAILABILITY

Source code and binaries are available at <<https://github.com/torognes/vsearch>>.

COPYRIGHT

Copyright (C) 2014-2025, Torbjørn Rognes, Frédéric Mahé and Tomás Flouri

All rights reserved.

Contact: Torbjørn Rognes <torognes@ifi.uio.no>, Department of Informatics, University of Oslo, PO Box

1080 Blindern, NO-0316 Oslo, Norway

This software is dual-licensed and available under a choice of one of two licenses, either under the terms of the GNU General Public License version 3 or the BSD 2-Clause License.

GNU General Public License version 3

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see ([link](#)) <<https://www.gnu.org/licenses/>> <<https://www.gnu.org/licenses/>>.

The BSD 2-Clause License

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

1. Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

We would like to thank the authors of the following projects for making their source code available:

- **vsearch** includes code from Google's CityHash project by Geoff Pike and Jyrki Alakuijala, providing some excellent hash functions available under a MIT license.
- **vsearch** includes code derived from Tatusov and Lipman's DUST program that is in the public domain.
- **vsearch** includes public domain code written by Alexander Peslyak for the MD5 message digest algorithm.
- **vsearch** includes public domain code written by Steve Reid and others for the SHA1 message digest algorithm.
- **vsearch** binaries may include code from the zlib library, copyright Jean-Loup Gailly and Mark Adler.
- **vsearch** binaries may include code from the bzip2 library, copyright Julian R. Seward.

SEE ALSO

swipe, an extremely fast pairwise local (Smith-Waterman) database search tool by Torbjørn Rognes, available at ([link](#)) <<https://github.com/torognes/swipe>> <<https://github.com/torognes/swipe>>.

swarm, a fast and accurate amplicon clustering method by Frédéric Mahé and Torbjørn Rognes, available at ([link](#)) <<https://github.com/torognes/swarm>> <<https://github.com/torognes/swarm>>.

VERSION HISTORY

New features and important modifications of **vsearch** (short lived or minor bug releases may not be mentioned):

v1.0.0 released November 28th, 2014

First public release.

v1.0.1 released December 1st, 2014

Bug fixes (sortbysize, semicolon after size annotation in headers) and minor changes (labels as secondary sort key for most sorts, treat T and U as identical for dereplication, only output size in --dbmatched file if --sizeout specified).

v1.0.2 released December 6th, 2014

Bug fixes (ssse3/sse4.1 requirement, memory leak).

v1.0.3 released December 6th, 2014

Bug fix (now writes help to stdout instead of stderr).

v1.0.4 released December 8th, 2014

Added --allpairs_global option. Reduce memory requirements slightly and eliminate memory leaks.

v1.0.5 released December 9th, 2014

Fixes a minor bug with --allpairs_global and --acceptall options.

v1.0.6 released December 14th, 2014

Fixes a memory allocation bug in chimera detection (--uchime_ref option).

v1.0.7 released December 19th, 2014

Fixes a bug in the output from chimera detection with the --uchimeout option.

v1.0.8 released January 22nd, 2015

Introduces several changes and bug fixes:

- a new linear memory aligner for alignment of sequences longer than 5,000 nucleotides,
- a new --cluster_size command that sorts sequences by decreasing abundance before clustering,
- meaning of userfields qlo, qhi, tlo, thi changed for compatibility with usearch,
- new userfields qilo, qih, tilo, tihi give alignment coordinates ignoring terminal gaps,
- in --uc output files, a perfect alignment is indicated with a '=' sign,
- the option --cluster_fast now sorts sequences by decreasing length, then by decreasing abundance and finally by sequence identifier,
- default --maxseqlength value set to 50,000 nucleotides,
- fix for bug in alignment in rare cases,
- fix for lack of detection of under- or overflow in SIMD aligner.

v1.0.9 released January 22nd, 2015

Fixes a bug in the function sorting sequences by decreasing abundance (--sortbysize).

v1.0.10 released January 23rd, 2015

Fixes a bug where the --sizein option was ignored and always treated as on, affecting clustering and dereplication commands.

v1.0.11 released February 5th, 2015

Introduces the possibility to output results in SAM format (for clustering, pairwise alignment and searching).

v1.0.12 released February 6th, 2015

Temporarily fixes a problem with long headers in FASTA files.

v1.0.13 released February 17th, 2015

Fix a memory allocation problem when computing multiple sequence alignments with the --msaout and --consout options, as well as a memory leak. Also increased line buffer for reading FASTA files to 4MB.

v1.0.14 released February 17th, 2015

Fix a bug where the multiple alignment and consensus sequence computed after clustering ignored the strand of the sequences. Also decreased size of line buffer for reading FASTA files to 1MB again due to excessive stack memory usage.

v1.0.15 released February 18th, 2015

Fix bug in calculation of identity metric between sequences when using the MBL definition (--id-def 3).

v1.0.16 released February 19th, 2015

Integrated patches from Debian for increased compatibility with various architectures.

v1.1.0 released February 20th, 2015

Added the --quiet option to suppress all output to stdout and stderr except for warnings and fatal errors. Added the --log option to write messages to a log file.

v1.1.1 released February 20th, 2015

Added info about --log and --quiet options to help text.

v1.1.2 released March 18th, 2015

Fix bug with large datasets. Fix format of help info.

v1.1.3 released March 18th, 2015

Fix more bugs with large datasets.

v1.2.0-1.2.19 released July 6th to September 8th, 2015

Several new commands and options added. Bugs fixed. Documentation updated.

v1.3.0 released September 9th, 2015

Changed to autotools build system.

v1.3.1 released September 14th, 2015

Several new commands and options. Bug fixes.

v1.3.2 released September 15th, 2015

Fixed memory leaks. Added '-h' shortcut for help. Removed extra 'v' in version number.

v1.3.3 released September 15th, 2015

Fixed bug in hexadecimal digits of MD5 and SHA1 digests. Added --samheader option.

v1.3.4 released September 16th, 2015

Fixed compilation problems with zlib and bzip2lib.

v1.3.5 released September 17th, 2015

Minor configuration/makefile changes to compile to native CPU and simplify makefile.

v1.4.0 released September 25th, 2015

Added --sizeorder option.

v1.4.1 released September 29th, 2015

Inserted public domain MD5 and SHA1 code to eliminate dependency on crypto and openssl libraries and their licensing issues.

v1.4.2 released October 2nd, 2015

Dynamic loading of libraries for reading gzip and bzip2 compressed files if available. Circumvention of missing gzoffset function in zlib 1.2.3 and earlier.

v1.4.3 released October 3rd, 2015

Fix a bug with determining amount of memory on some versions of Apple OS X.

v1.4.4 released October 3rd, 2015

Remove debug message.

v1.4.5 released October 6th, 2015

Fix memory allocation bug when reading long FASTA sequences.

v1.4.6 released October 6th, 2015

Fix subtle bug in SIMD alignment code that reduced accuracy.

v1.4.7 released October 7th, 2015

Fixes a problem with searching for or clustering sequences with repeats. In this new version, vsearch looks at all words occurring at least once in the sequences in the initial step. Previously only words occurring exactly once were considered. In addition, vsearch now requires at least 10 words to be shared by the sequences, previously only 6 were required. If the query contains less than 10 words, all words must be present for a match. This change seems to lead to slightly reduced recall, but somewhat increased precision, ending up with slightly improved overall accuracy.

v1.5.0 released October 7th, 2015

This version introduces the new option --minwordmatches that allows the user to specify the minimum number of matching unique words before a sequence is considered further. New default values for different word lengths are also set. The minimum word length is increased to 7.

v1.6.0 released October 9th, 2015

This version adds the relabeling options (--relabel, --relabel_md5 and --relabel_sha1) to the shuffle command. It also adds the --xsize option to the clustering, dereplication, shuffling and sorting commands.

v1.6.1 released October 14th, 2015

Fix bugs and update manual and help text regarding relabelling. Add all relabelling options to the subsampling command. Add the --xsize option to chimera detection, dereplication and fastq filtering commands. Refactoring of code.

v1.7.0 released October 14th, 2015

Add --relabel_keep option.

v1.8.0 released October 19th, 2015

Added --search_exact, --fastx_mask and --fastq_convert commands. Changed most commands to read FASTQ input files as well as FASTA files. Modified --fastx_revcomp and --fastx_subsample to write FASTQ files.

v1.8.1 released November 2nd, 2015

Fixes for compatibility with QIIME and older OS X versions.

v1.9.0 released November 12th, 2015

Added the --fastq_mergepairs command and associated options. This command has not been tested well yet. Included additional files to avoid dependency of autoconf for compilation. Fixed an error where identifiers in fasta headers were not truncated at tabs, just spaces. Fixed a bug in detection of the file format (FASTA/FASTQ) of a gzip compressed input file.

v1.9.1 released November 13th, 2015

Fixed memory leak and a bug in score computation in --fastq_mergepairs, and improved speed.

v1.9.2 released November 17th, 2015

Fixed a bug in the computation of some values with --fastq_stats.

v1.9.3 released November 19th, 2015

Workaround for missing x86intrin.h with old compilers.

v1.9.4 released December 3rd, 2015

Fixed incrementation of counter when relabeling dereplicated sequences.

v1.9.5 released December 3rd, 2015

Fixed bug resulting in inferior chimera detection performance.

v1.9.6 released January 8th, 2016

Fixed bug in aligned sequences produced with --fastapairs and --userout (qrow, trow) options.

v1.9.7 released January 12th, 2016

Masking behaviour is changed somewhat to keep the letter case of the input sequences unchanged when no masking is performed. Masking is now performed also during chimera detection. Documentation updated.

v1.9.8 released January 22nd, 2016

Fixed bug causing segfault when chimera detection is performed on extremely short sequences.

v1.9.9 released January 22nd, 2016

Adjusted default minimum number of word matches during searches for improved performance.

v1.9.10 released January 25th, 2016

Fixed bug related to masking and lower case database sequences.

v1.10.0 released February 11th, 2016

Parallelized and improved merging of paired-end reads and adjusted some defaults. Removed progress indicator when stderr is not a terminal. Added --fasta_score option to report chimera scores in FASTA files. Added --rereplicate and --fastq_eestats commands. Fixed typos. Added re-labelling to files produced with --consout and --profile options.

v1.10.1 released February 23rd, 2016

Fixed a bug affecting the --fastq_mergepairs command causing FASTQ headers to be truncated at first space (despite the bug fix release 1.9.0 of November 12th, 2015). Full headers are now included in the output (no matter if --notrunclabels is in effect or not).

v1.10.2 released March 18th, 2016

Fixed a bug causing a segmentation fault when running --usearch_global with an empty query sequence. Also fixed a bug causing imperfect alignments to be reported with an alignment string of '=' in uc output files. Fixed typos in man file. Fixed fasta/fastq processing code regarding presence or absence of compression library header files.

v1.11.1 released April 13th, 2016

Added strand information in UC file for --derep_fulllength and --derep_prefix. Added expected errors (ee) to header of FASTA files specified with --fastaout and --fastaout_discarded when --eeout or --fastq_eeout option is in effect for fastq_filter and fastq_mergepairs. The options --eeout and --fastq_eeout are now equivalent.

v1.11.2 released June 21st, 2016

Two bugs were fixed. The first issue was related to the --query_cov option that used a different coverage definition than the qcov userfield. The coverage is now defined as the fraction of the whole query sequence length that is aligned with matching or mismatching residues in the target. All gaps are ignored. The other issue was related to the consensus sequences produced during clustering when only N's were present in some positions. Previously these would be converted to A's in the consensus. The behaviour is changed so that N's are produced in the consensus, and it should now be more compatible with usearch.

v2.0.0 released June 24th, 2016

This major new version supports reading from pipes. Two new options are added: --gzip_decompress and --bzip2_decompress. One of these options must be specified if reading compressed input from a pipe, but are not required when reading from ordinary files. The vsearch header that was previously written to stdout is now written to stderr. This enables piping of results for further processing. The file name '-' now represent standard input (/dev/stdin) or standard output (/dev/stdout) when reading or writing files, respectively. Code for reading FASTA and FASTQ files has been refactored.

v2.0.1 released June 30th, 2016

Avoid segmentation fault when masking very long sequences.

v2.0.2 released July 5th, 2016

Avoid warnings when compiling with GCC 6.

v2.0.3 released August 2nd, 2016

Fixed bad compiler options resulting in Illegal instruction errors when running precompiled binaries.

v2.0.4 released September 1st, 2016

Improved error message for bad FASTQ quality values. Improved manual.

v2.0.5 released September 9th, 2016

Add options --fastaout_discarded and --fastqout_discarded to output discarded sequences from subsampling to separate files. Updated manual.

v2.1.0 released September 16th, 2016

New command: --fastx_filter. New options: --fastq_maxlen, --fastq_truncee. Allow --minwordmatches down to 3.

v2.1.1 released September 23rd, 2016

Fixed bugs in output to UC-files. Improved help text and manual.

v2.1.2 released September 28th, 2016

Fixed incorrect abundance output from fastx_filter and fastq_filter when relabelling.

v2.2.0 released October 7th, 2016

Added OTU table generation options --biomout, --mothur_shared_out and --otutabout to the clustering and searching commands.

v2.3.0 released October 10th, 2016

Allowed zero-length sequences in FASTA and FASTQ files. Added --fastq_trunclen_keep option. Fixed bug with output of OTU tables to pipes.

v2.3.1 released November 16th, 2016

Fixed bug where --minwordmatches 0 was interpreted as the default minimum word matches for the given word length instead of zero. When used in combination with --maxaccepts 0 and --maxrejects 0 it will allow complete bypass of kmer-based heuristics.

v2.3.2 released November 18th, 2016

Fixed bug where vsearch reported the ordinal number of the target sequence instead of the cluster number in column 2 on H-lines in the uc output file after clustering. For search and alignment commands both usearch and vsearch reports the target sequence number here.

v2.3.3 released December 5th, 2016

A minor speed improvement.

v2.3.4 released December 9th, 2016

Fixed bug in output of sequence profiles and updated documentation.

v2.4.0 released February 8th, 2017

Added support for Linux on Power8 systems (ppc64le) and Windows on x86_64. Improved detection of pipes when reading FASTA and FASTQ files. Corrected option for specifying output from fastq_eestats command in help text.

v2.4.1 released March 1st, 2017

Fixed an overflow bug in fastq_stats and fastq_eestats affecting analysis of very large FASTQ files.

Fixed maximum memory usage reporting on Windows.

v2.4.2 released March 10th, 2017

Default value for fastq_minovlen increased to 16 in accordance with help text and for compatibility with usearch. Minor changes for improved accuracy of paired-end read merging.

v2.4.3 released April 6th, 2017

Fixed bug with progress bar for shuffling. Fixed missing N-lines in UC files with usearch_global, search_exact and allpairs_global when the output_no_hits option was not specified.

v2.4.4 released August 28th, 2017

Fixed a few minor bugs, improved error messages and updated documentation.

v2.5.0 released October 5th, 2017

Support for UDB database files. New commands: fastq_stripright, fastq_eestats2, makeudb_usearch, udb2fasta, udbinfo, and udbstats. New general option: no_progress. New options minsize and maxsize to fastx_filter. Minor bug fixes, error message improvements and documentation updates.

v2.5.1 released October 25th, 2017

Fixed bug with bad default value of 1 instead of 32 for minseqlength when using the makeudb_usearch command.

v2.5.2 released October 30th, 2017

Fixed bug with where '-' as an argument to the fastq_eestats2 option was treated literally instead of equivalent to stdin.

v2.6.0 released November 10th, 2017

Rewritten paired-end reads merger with improved accuracy. Decreased default value for fastq_minnovlen option from 16 to 10. The default value for the fastq_maxdifs option is increased from 5 to 10. There are now other more important restrictions that will avoid merging reads that cannot be reliably aligned.

v2.6.1 released December 8th, 2017

Improved parallelisation of paired end reads merging.

v2.6.2 released December 18th, 2017

Fixed option xsize that was partially inactive for commands uchime_denovo, uchime_ref, and fastx_filter.

v2.7.0 released February 13th, 2018

Added commands cluster_unoise, uchime2_denovo and uchime3_denovo contributed by Davide Albanese based on Robert Edgar's papers. Refactored fasta and fastq print functions as well as code for extraction of abundance and other attributes from the headers.

v2.7.1 released February 16th, 2018

Fix several bugs on Windows related to large files, use of "-" as a file name to mean stdin or stdout, alignment errors, missed kmers and corrupted UDB files. Added documentation of UDB-related commands.

v2.7.2 released April 20th, 2018

Added the sintax command for taxonomic classification. Fixed a bug with incorrect FASTA headers of consensus sequences after clustering.

v2.8.0 released April 24th, 2018

Added the fastq_maxdiffpct option to the fastq_mergepairs command.

v2.8.1 released June 22nd, 2018

Fixes for compilation warnings with GCC 8.

v2.8.2 released August 21st, 2018

Fix for wrong placement of semicolons in header lines in some cases when using the sizeout or xsize options. Reduced memory requirements for full-length dereplication in cases with many duplicate sequences. Improved wording of fastq_mergepairs report. Updated manual regarding use of sizein and sizeout with dereplication. Changed a compiler option.

v2.8.3 released August 31st, 2018

Fix for segmentation fault for --derep_fulllength with --uc.

v2.8.4 released September 3rd, 2018

Further reduce memory requirements for dereplication when not using the uc option. Fix output during subsampling when quiet or log options are in effect.

v2.8.5 released September 26th, 2018

Fixed a bug in fastq_eestats2 that caused the values for large lengths to be much too high when the input sequences had varying lengths.

v2.8.6 released October 9th, 2018

Fixed a bug introduced in version 2.8.2 that caused derep_fulllength to include the full FASTA header in its output instead of stopping at the first space (unless the notrunclabels option is in effect).

v2.9.0 released October 10th, 2018

Added the fastq_join command.

v2.9.1 released October 29th, 2018

Changed compiler options that select the target cpu and tuning to allow the software to run on any 64-bit x86 system, while tuning for more modern variants. Avoid illegal instruction error on some architectures. Update documentation of rereplicate command.

v2.10.0 released December 6th, 2018

Added the sff_convert command to convert SFF files to FASTQ. Added some additional option argument checks. Fixed segmentation fault bug after some fatal errors when a log file was specified.

v2.10.1 released December 7th, 2018

Improved sff_convert command. It will now read several variants of the SFF format. It is also able to read from a pipe. Warnings are given if there are minor problems. Errors messages have been improved. Minor speed and memory usage improvements.

v2.10.2 released December 10th, 2018

Fixed bug in syntax with reversed order of domain and kingdom.

v2.10.3 released December 19th, 2018

Ported to Linux on ARMv8 (aarch64). Fixed compilation warning with gcc version 8.1.0 and 8.2.0.

v2.10.4 released January 4th, 2019

Fixed serious bug in x86_64 SIMD alignment code introduced in version 2.10.3. Added link to BioConda in README. Fixed bug in fastq_stats with sequence length 1. Fixed use of equals symbol in UC files for identical sequences with cluster_fast.

v2.11.0 released February 13th, 2019

Added ability to trim and filter paired-end reads using the reverse option with the fastx_filter and fastq_filter commands. Added --xee option to remove ee attributes from FASTA headers. Minor invisible improvement to the progress indicator.

v2.11.1 released February 28th, 2019

Minor change to the handling of the weak_id and id options when using cluster_unequal.

v2.12.0 released March 19th, 2019

Take sequence abundance into account when computing consensus sequences or profiles after clustering. Warn when rereplicating sequences without abundance info. Guess offset 33 in more cases with fastq_chars. Stricter checking of option arguments and option combinations.

v2.13.0 released April 11th, 2019

Added the --fastx_getseq, --fastx_getsegs and --fastx_getsubseq commands to extract sequences from a FASTA or FASTQ file based on their labels. Improved handling of ambiguous nucleotide symbols. Corrected behaviour of --uchime_ref command with and options --self and --selfid. Strict detection of illegal options for each command.

v2.13.1 released April 26th, 2019

Minor changes to the allowed options for each command. All commands now allow the log, quiet and threads options. If more than 1 thread is specified for commands that are not multi-threaded, a warning will be issued. Minor changes to the manual.

v2.13.2 released April 30th, 2019

Fixed bug related to improper handling of newlines on Windows. Allowed option strand plus to uchime_ref for compatibility.

v2.13.3 released April 30th, 2019

Fixed bug in FASTQ parsing introduced in version 2.13.2.

v2.13.4 released May 10th, 2019

Added information about support for gzip- and bzip2-compressed input files to the output of the version command. Adapted source code for compilation on FreeBSD and NetBSD systems.

v2.13.5 released July 2nd, 2019

Added cut command to fragment sequences at restriction sites. Silenced output from the fastq_stats command if quiet option was given. Updated manual.

v2.13.6 released July 2nd, 2019

Added info about cut command to output of help command.

v2.13.7 released September 2nd, 2019

Fixed bug in consensus sequence introduced in version 2.13.0.

v2.14.0 released September 11th, 2019

Added relabel_self option. Made fasta_width, sizein, sizeout and relabelling options valid for certain commands.

v2.14.1 released September 18th, 2019

Fixed bug with sequences written to file specified with fastaout_rev for commands fastx_filter and fastq_filter.

v2.14.2 released January 28th, 2020

Fixed some issues with the cut, fastx_revcomp, fastq_convert, fastq_mergepairs, and makefdb_usearch commands. Updated manual.

v2.15.0 released June 19th, 2020

Update manual and documentation. Turn on notrunclabels option for syntax command by default. Change maxhits 0 to mean unlimited hits, like the default. Allow non-ascii characters in headers, with a warning. Sort centroids and uc too when clusterout_sort specified. Add cluster id to centroids output when clusterout_id specified. Improve error messages when parsing FASTQ files. Add missing fastq_qminout option and fix label_suffix option for fastq_mergepairs. Add derep_id command that dereplicates based on both label and sequence. Remove compilation warnings.

v2.15.1 released October 28th, 2020

Fix for dereplication when including reverse complement sequences and headers. Make some extra checks when loading compression libraries and add more diagnostic output about them to the output of the version command. Report an error when fastx_filter is used with FASTA input and options that require FASTQ input. Update manual.

v2.15.2 released January 26th, 2021

No real functional changes, but some code and compilation changes. Compiles successfully on macOS running on Apple Silicon (ARMv8). Binaries available. Code updated for C++11. Minor adaptations for Windows compatibility, including the use of the C++ standard library for regular expressions. Minor changes for compatibility with Power8. Switch to C++ header files.

v2.16.0 released March 22nd, 2021

This version adds the orient command. It also handles empty input files properly. Documentation has been updated.

v2.17.0 released March 29nd, 2021

The fastq_mergepairs command has been changed. It now allows merging of sequences with overlaps as short as 5 bp if the --fastq_minovlen option has been adjusted down from the default 10. In addition, much fewer pairs of reads should now be rejected with the reason 'multiple potential alignments' as the algorithm for detecting those have been changed.

v2.17.1 released June 14th, 2021

Modernized code. Minor changes to help info.

v2.18.0 released August 27th, 2021

Added the fasta2fastq command. Fixed search bug on ppc64le. Fixed bug with removal of size and ee info in uc files. Fixed compilation errors in some cases. Made some general code improvements. Updated manual.

v2.19.0 released December 21st, 2021

Added the lcaout and lca_cutoff options to enable the output of last common ancestor (LCA) information about hits when searching. The randseed option was added as a valid option to the sintax command. Code improvements.

v2.20.0 released January 10th, 2022

Added the fastx_uniques command and the fastq_qout_max option for dereplication of FASTQ files. Some code cleaning.

v2.20.1 released January 11th, 2022

Fixes a bug in fastq_mergelpair that caused an occational hang at the end when using multiple threads.

v2.21.0 released January 12th, 2022

This version adds the sample, qsegout and tsegout options. It enables the use of UDB databases with uchime_ref.

v2.21.1 released January 18th, 2022

Fix a problem with dereplication of empty input files. Update Altivec code on ppc64le for improved compiler compatibility (vector->__vector).

v2.21.2 released September 12th, 2022

Fix problems with the lcaout option when using maxaccepts above 1 and either lca_cutoff below 1 or with top_hits_only enabled. Update documentation. Update code to avoid compiler warnings.

v2.22.0 released September 19th, 2022

Add the derep_smallmem command for dereplication using little memory.

v2.22.1 released September 19th, 2022

Fix compiler warning.

v2.23.0 released July 7th, 2023

Update documentation. Add citation file. Modernize and improve code. Fix several minor bugs. Fix compilation with GCC 13. Print stats after fastq_mergepairs to log file instead of stderr. Handle sizein option correctly with dbmatched option for usearch_global. Allow maxseqlength option for makeudb_usearch. Fix memory allocation problem with chimera detection. Add lengthout and xlength options. Increase precision for eeout option. Add warning about sintax algorithm, random seed and multiple threads. Refactor chimera detection code. Add undocumented experimental long_chimeras_denovo command. Fix segfault with clustering. Add more references.

v2.24.0 released October 26th, 2023

Update documentation. Improve code. Allow up to 20 parents for the undocumented and experimental chimeras_denovo command. Fix compilation warnings for sha1.c. Compile for release (not debug) by default.

v2.25.0 released November 10th, 2023

Allow a given percentage of mismatches between chimeras and parents for the experimental chimeras_denovo command.

v2.26.0 released November 24th, 2023

Enable the maxseqlength and minseqlength options for the chimera detection commands. When the usearch_global or search_exact commands are used, OTU tables will include samples and OTUs with no matches.

v2.26.1 released November 25th, 2023

No real changes, but the previous version was released without proper updates to the source code.

v2.27.0 released January 19th, 2024

The usearch_global and search_exact commands now support FASTQ files as well as FASTA files as input. This version of vsearch includes clarifications and updates to the manual. Some code has been refactored. Generic Dockerfiles for major Linux distributions have been included. Some warnings from compilers and other tools have been eliminated. The release for Windows will also include DLL's for the two compression libraries.

v2.27.1 released April 6th, 2024

This version fixes the weak_id option and makes searches report weak hits in some cases. It also updates the names of the compression libraries to libz.so.1 and libbz2.so.1 on Linux to make them work on common Linux distributions without installing additional packages. README.md has been updated with information about compression libraries on Windows.

v2.28.0 released April 26th, 2024

The sintax command has been improved in several ways in this version of vsearch. Please note that several details of this algorithm is not clearly described in the preprint, and the implementation in vsearch differs from that in usearch. The former vsearch version did not always choose the most common taxonomic entity over the 100 bootstraps among the database sequences with the highest amount of word similarity to the query. Instead, if several sequences had an equal similarity with the query, the sequence encountered in the earliest bootstrap was chosen. The confidence level was calculated based on this sequence compared to the selected sequences from the other 99 bootstraps. This could lead to a suboptimal choice with a low confidence. In the new version, the most common of the sequences with the highest amount of word similarity across the 100 bootstraps will be selected, and ties will be broken randomly. Another problem with the old implementation was that if several sequences had the same amount of word similarity, the shortest one in the reference database would be chosen, and if they were equally long, the earliest in the database file would be chosen. A new option called sintax_random has now been introduced. This option will randomly select one of the sequences with the highest number of shared words with the query, without considering their length or position. This avoids a bias towards shorter reference sequences. This option is strongly recommended and will probably soon be the default. Furthermore, a ninth taxonomic rank, strain (letter t), is now recognized. The speed of the sintax command has also been significantly improved at least in some cases. Run vsearch with the randseed option and 1 thread to ensure reproducibility of the random choices in the algorithm.

v2.28.1 released April 26th, 2024

Fix a segmentation fault that could occur with the blast6out and output_no_hits options.

v2.29.0 released September 26th, 2024

This version fixes seven bugs (see changelog below), adds initial support for RISC-V architectures, and improves code quality and code testing (1,210 new tests):

- add: experimental support for RISCV64 and other 64-bit little-endian architectures, thanks to Michael R. Crusoe and his fellow Debian developers (issue #566),
- add: official support for clang-19 and gcc 14,
- add: beta support for clang-20,
- remove: unused --output option for command --fastq_stats (issue #572),
- fix: bug in --sintax when selecting the best lineage (only low confidence values below 0.5 were affected) (issue #573),

- fix: out-of-bounds error in --fastq_stats when processing empty reads (issue #571),
- fix: bug in --cut, patterns with multiple cutting sites were not detected (commit 4c4f9fa70f14b28d50185dbf322cf5727087e86a),
- fix: memory error (segmentation fault) when using --derep_id and --strand (issue #565),
- fix: --fastq_join now obeys to --quiet and --log options (commit 87f968b09f17c17ebf8db00aebe86e89b13a3948),
- fix: --fastq_join quality padding is now also set to Q40 when quality offset is 64 (commit be0bf9b48d782286c4ce38f0bf1a4c82bd230250),
- fix: (partial) --fastq_join's handling of abundance annotations (commit f2bbcb421dc2f4dfa6603b9f31ec3e4598c1b591),
- improve: additional safeguards to validate input values and to make sure that they are within acceptable limits. Changes concern options --abskew (commit a530dd8990f8a05cb25fc0b6a5da5a14d28fbedd) and --fastq_maxdiffs (commit 4b254db7f120bfd49e86185ef3cd9070c236f940),
- improve: code quality (1.3k+ commits, 6k+ clang-tidy warnings eliminated),
- improve: documentation and help messages (issue #568),
- improve: complete refactoring and modernization of a subset of commands (--sortbylength, --sortysize, --shuffle, --rereplicate, --cut, --fastq_join, --fasta2fastq, --fastq_chars),
- improve: code-coverage of our test-suite for the above-mentioned commands (1,210 new tests, 4,753 in total)

v2.29.1 released October 24th, 2024

Fix a segmentation fault that could occur during alignment in version 2.29.0, for example with --uchime_ref. Some improvements to code and documentation.

v2.29.2 released December 20th, 2024

Fix a segmentation fault during clustering when the set of clusters is empty. Initial documentation in markdown format available on GitHub Pages.

v2.29.3 released February 3rd, 2025

This version is released in order to mitigate a bug that occurs when compiling the ‘align_simd.cc’ file on x86_64 systems with the GNU C++ compiler version 9 or later with the ‘-O3’ optimization option. It results in incorrect code that may cause bad alignments in some circumstances. We are investigating this issue further, but for now we recommend compiling with the ‘-O2’ flag. The README.md file and the Dockerfiles have been updated to reflect this. The binaries released with this version will include this fix.

v2.29.4 released February 14th, 2025

Adjust the window size used for chimera detection down from 64 to 32. The window size was by accident increased from 32 to 64 in version 2.23.0, leading to somewhat fewer chimeras being predicted. In addition, a compiler pragma has been included in align_simd.cc to further protect the compiler from generating wrong code.

v2.30.0 released February 27th, 2025

Add options ‘--n_mismatch’, ‘--fastq_minqual’, and ‘--fastq_truncee_rate’. The ‘--n_mismatch’ option will count N’s as mismatches in alignments, which may be useful to get sensible alignments for sequences with lots of N’s. By default N’s are counted as matches. Both the scoring and the counting of matches are affected. The new ‘--fastq_minqual’ option for the ‘fastq_filter’ and ‘fastx_filter’ commands will discard sequences with any bases with a quality scores below the given value. The new ‘--fastq_truncee_rate’ option for the same commands will truncate sequences at the first position where the number of expected errors per base is above the given value.

v2.30.1 released October 3rd, 2025

This version incorporates many code improvements, more extensive testing, better documentation and some minor bug fixes. List of changes:

- fix: use-after-free introduced in commit de6c1d8 (Jun 13, 2024),
- fix: (harmless) out-of-bounds memory issue in --derep_prefix (commit 8a0a508b),
- fix: (harmless) memory leak in --fastx_getseqs --label_field (commit a9c42713),
- fix: (harmless) memory leak when using option --userfields (--allpairs_global commit 03b95bcf; --cluster_* commit 2fde5472; --search_exact commit 45cd56d6; --usearch_global commit d83bfee9),
- fix: (harmless) valgrind error, use of uninitialized values (commit 8bab2444), also eliminates a pesky compilation warning,
- change: passing a negative value to --fastq_truncee_rate is now an error (commit a120f371),
- change: passing a negative value to --fastq_minqual is now an error (commit ff5b0c99),
- change: passing a non-ASCII symbol to --join_padgap or --join_padgapq is now an error (commit a708f5b3),
- change: when using --gapopen "*" to forbid gap opening, the penalty is now set to INT_MAX (rather than 1,000). This might change alignment results for users who relied on the old behavior (thanks to Denis Filloux, issue #602, commit 96e9cf9e),
- change: when using command --chimeras_denovo, --tabbedout or --alnout can be the only output files specified (commit d51f0a456300a0ea69c035ba3de23c5ecf3da348)
- change: when using command --chimeras_denovo, option --lengthout is now accepted (commit 3f55cc6b)
- change: when using command --chimeras_denovo, option --xlength is now accepted (commit 0fd346cd)
- add: compilation option GLIBCXX_DEBUG when compiling for debugging (commit 5cf4a6c1, option activates more runtime checks),
- add: official support for clang 20,
- add: initial support for clang 21, initial support for GCC 15,
- add: experimental support for clang 22,
- improve: more accurate line number when reporting illegal characters in fastq headers (commit 539084e9),
- improve: more accurate line number when reporting non-ASCII characters in fastq headers (commit 98a851ed),
- improve: remove checks for unneeded libraries during compilation (commit 249bb5d5276b1be6c9add60a538a16e1b9d0ebfb),
- improve: code quality (8,536 clang-tidy warnings eliminated),
- improve: documentation and help messages (issues #604),
- improve: complete refactoring and modernization of the command --fastq_stats,
- improve: command --fastq_stats is now up to twice faster (tested on x86-64),
- improve: extensive test-suites for --fastq_stats and --sff_convert,
- improve: code coverage of our test-suite